

IS THERE A PLATONIC TREE?

CALIBRATING LATENT HYPERBOLICITY

IN FOUNDATION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

Hyperbolic methods for representation learning rest on a premise we call latent hyperbolicity: standard models are already tree-like, because their class geometry scores a low Gromov δ . We show that this raw reading is confounded by dimension, by covariance spectrum and by the supremum statistic itself. We build an instrument that reads every score as an excess over a random cloud with the same dimension and spectrum, ranks that excess against matched null replicates, and adds a depth test whose power and validity regime are measured. Applied to 12 vision backbones, 6 datasets and 16 text models, the instrument gives a nuanced answer. On image features, where the premise is read, the excess sits within null noise in most cells and is small where it survives. On class centroids, clustered structure is genuine in 49 of 72 cells, but it is rarely deeper than a star. The depth test, conservative and weak at ImageNet’s noise level, certifies hierarchy above the superclasses in 4 of 12 ImageNet backbones, and a backbone trained in hyperbolic space shows the same clustering and no detected depth. The trees are moderately shared: a naive comparison isolates the self-supervised family as an island, an artifact of the clustering cut, and once the cut is controlled every training recipe recovers the human taxonomy partially, more so with supervision, with the self-supervised tree living in angles rather than distances. Read correctly, foundation models organize classes into clustered structure that is occasionally hierarchical and moderately shared; they do not converge to one common tree, and their raw tree-likeness is not evidence for hyperbolic geometry.

1 INTRODUCTION

Is there a Platonic tree? Read correctly, foundation models organize classes into clustered structure that is occasionally hierarchical and moderately shared; they do not converge to one common tree, and their raw tree-likeness is not evidence for hyperbolic geometry.

figures/fig1_concept.pdf (conceptual figure, drawn by the authors; same-size placeholder)

Figure 1: **A random cloud, a star and a tree all read as tree-like; only the instrument tells them apart.** Sketched: a high-dimensional random cloud, clusters around a hub and a nested hierarchy, all with a low raw reading, the excess that separates the cloud from the other two, and the depth test that separates the star from the tree.

Semantic knowledge is hierarchical, and hyperbolic space is the natural geometric home for hierarchies (Nickel & Kiela, 2017; Sarkar, 2011). A line of work therefore imposes it on representations at training time, from vision to vision–language and language (Khrulkov et al., 2020; Ermolov et al., 2022; Atigh et al., 2022; Bdeir et al., 2024; Desai et al., 2023; Pal et al., 2025; Chen et al., 2022; He et al., 2025; Sinha et al., 2024). These methods cite a premise, which we call *latent hyperbolicity*: standard models are said to organize concepts already in a tree-like metric structure, because the Gromov δ of their features is low (Khrulkov et al., 2020; Bdeir et al., 2024; Yang et al., 2025; Tifrea et al., 2019). Is the premise right, and what does it license?

The raw reading is confounded. Gromov δ assigns zero to trees (Gromov, 1987), but three artifacts push it down without any hierarchy. High-dimensional clouds concentrate toward equidistance, which is a degenerate star tree: we call this the *dimension confound*. Anisotropic spectra mimic low-dimensional behavior: the *spectrum confound*. And the supremum over sampled quadruples is a one-quadruple extreme that does not converge: the *statistic confound*. The three are visible in real models: small GPT-2 models sit at their matched null despite a low raw reading, and classic sentence embedders share their raw reading with vision-contrastive models while carrying none of their structure.

We build the *instrument* that reads δ correctly. It calibrates the estimator on reference geometries, expresses every reading as an *excess* over a null cloud matched in dimension and covariance spectrum, ranks that excess against matched replicates, and adds a *depth test* whose power and validity regime are measured. We apply it at the sample level, where the premise is read, and at the class level, where hierarchy lives. With it we ask of geometry what convergence work asks of similarity (Huh et al., 2024; Gröger et al., 2026): what survives calibration, and is it shared? In the words of the title: is there a Platonic tree? They calibrate similarity across models; we calibrate geometry within a model, and comparing trees across models needs a calibration of its own.

The answer has three parts. On image features the excess sits within null noise in most cells and is small where it survives, so the premise does not survive calibration as stated. On class centroids, clustered structure is genuine in 49 of 72 cells, but it is rarely deeper than a star. The depth test, conservative and weak at ImageNet’s noise level, certifies hierarchy above the superclasses in 4 of 12 ImageNet backbones, and a backbone trained in hyperbolic space shows the same clustering and no detected depth. The trees are moderately shared: the naive comparison manufactures an island, and once the clustering cut is controlled every recipe recovers the human taxonomy partially, more so with supervision, with the self-supervised structure living in angles. What follows for practice is what to measure before imposing curvature and which zero-cost readout collects what is there. Clustered, occasionally hierarchical, moderately shared: not one common tree, and no license for curvature.

Our contributions are:

1. **The instrument.** Estimator calibration, spectrum-matched nulls read as corrected ranks, and a depth test with measured power; the raw reading misleads at the sample level and at the class level alike.
2. **The census.** Clustered structure in most models, hierarchy certified in a few, and a hyperbolic backbone as the control for imposing the geometry.
3. **The map of trees.** The island is an artifact of the cut; sharing is graded and supervision-dependent, and the self-supervised tree is angular.
4. **Consequences for practice.** What to measure before imposing curvature, and which zero-cost readout collects what is there.

The premise and its confounds have a history, which we review before building the instrument.

2 RELATED WORK AND BACKGROUND

The premise. Hyperbolic representation learning builds on tree embeddability and Poincaré embeddings and now spans vision, vision–language, graphs and language (Sarkar, 2011; Nickel & Kiela, 2017; Ganea et al., 2018; Chami et al., 2019). Its stated premise rests on a low raw reading for CNN and ViT features, token embeddings and word vectors, taken at the sample level, uncali-




Figure 2: **The raw reading cannot be read on its own, and the naive tree comparison manufactures the island.** Plotted: (a) the raw reading of the 12 backbones on ImageNet against their matched null means, shown hollow; (b) a sample-level cell (ViT-T on DTD images) and a class-level cell (SigLIP-B on CIFAR-10 centroids) with the same raw reading and opposite verdicts; (c) DINOv2-vs-block tree agreement before and after the configuration control of Section 5.

brated and under a single construction of the statistic (Khrulkov et al., 2020; Bdeir et al., 2024; Yang et al., 2025; Tifrea et al., 2019). Section 4 re-reads them with the instrument.

Calibrating geometry against nulls. Network science reads Gromov δ on real graphs against reference graphs: the curvature of communication and social networks is compared with that of lattices, hyperbolic grids and Erdős–Rényi random graphs of comparable size, and δ is read relative to the graph’s diameter (Narayan & Sanjeev, 2011; Kennedy et al., 2013). Adcock et al. (2013) pair δ with tree decompositions and find the bare value of δ too coarse a guide to tree-likeness on its own. Reading δ against a reference is therefore not new. Representation similarity, intrinsic dimension and neural collapse characterize Euclidean structure without testing tree-likeness (Kornblith et al., 2019; Ansuini et al., 2019; Pope et al., 2021; Papayan et al., 2020). The Platonic Representation Hypothesis argues that models converge toward shared representations; Gröger et al. (2026) show that permutation-calibrated similarity removes most of that convergence while local structure survives, and Park et al. (2024) read concept hierarchy in language models from orthogonality (Huh et al., 2024; Koepke et al., 2026). What is new here is the instrument for embedding clouds rather than graphs: spectrum-matched nulls at the resolution of the null replicate, a depth test against matched stars measured on real clouds, a hyperbolic backbone as the control for imposing the geometry, and a calibration of the tree-to-tree comparison itself.

Background. For four points w, x, y, z of a metric space, sort the three pairing sums $d(w, x) + d(y, z)$, $d(w, y) + d(x, z)$, $d(w, z) + d(x, y)$ as $S_{\max} \geq S_{\text{mid}} \geq S_{\min}$; the four-point defect is $(S_{\max} - S_{\text{mid}})/2$ and δ is its supremum. Trees have $\delta = 0$, and under this definition the hyperbolic plane has $\delta = \ln 2 \approx 0.693$, distinct from the thin-triangles constant $\ln(1 + \sqrt{2}) \approx 0.881$ (Nica & Špakula, 2016). We use the normalized $\delta_{\text{norm}} = \delta_{\max}/\text{diam}$ throughout, on which trees $<$ hyperbolic $<$ spherical. Sampled-quadruple estimation has known bias and scale caveats, which the calibration-plus-null design addresses, and Ollivier–Ricci curvature on a kNN graph is the local companion (Fournier et al., 2015; Ollivier, 2009). How the raw reading becomes a calibrated one is the subject of the next section.

3 THE INSTRUMENT

If the premise were right, a low raw reading would be evidence on its own. The dimension confound says it is not, so this section builds the measurement layer in whose units every later claim is stated.

The premise lives at the sample level; hierarchy lives at the class level. Each frozen model is a fixed feature extractor; the premise is read on per-image features, where the literature measures it. Hierarchy is measured on class centroids: in vision, the mean of each class’s cached training embeddings, a hundred per class on ImageNet; in text, the embeddings of “a photo of a {class}” over the ImageNet classes. A *cell* is one pair of model and class set; the census covers 12 vision backbones on 6 datasets and 16 text models. The raw reading is the 99.9th-percentile four-point statistic $\hat{\delta}_{99.9}$ over 5×10^5 sampled quadruples per seed, averaged over 10 seeds; the supremum,




Figure 3: **Clustered structure is the rule at the class level.** Plotted: (a) the excess of every cell of the census, with hollow markers for cells that are not genuine; (b) the excess against model scale for each dataset, where only DINOv2 on CIFAR-10 deepens monotonically.

the pre-registered statistic, is reported in the appendix for the reason given below. The excess is budget-stable (Appendix A.15).

A low raw reading is not evidence. Calibrated on reference geometries, the estimator recovers the tree zero and the $\ln 2$ constant of $\mathbb{H}^2$, and spheres stay high. Iid Gaussians, which have no structure at all, drift lower as the dimension grows, from 0.104 at $d=192$ to 0.046 at $d=1536$. That drift is the dimension confound in its purest form: the raw reading of a real backbone must be compared with the reading of a structureless cloud of the same shape (Table 3). The next paragraph builds that comparison.

The reading is an excess over a matched null. For every cell we synthesize null clouds with the same n , d , mean and covariance spectrum. We use the Haar construction throughout: random orthogonal coefficients recombined with the real singular values, so the sample spectrum is exact; the Gaussian-coefficient construction is the alternative reported in the appendix. In plain terms, the spectrum null is a cloud with the same shape as the real one and no structure inside it; the hub null of Section 4 keeps every cluster of the real cloud and only rearranges where the clusters sit. The **excess** is the raw reading minus the mean reading of 200 such replicates. The primary evidence is the *percentile rank*, the number of replicates above the real value, and its left-tail p . A cell is *genuine* when its Benjamini–Hochberg-corrected p across the census is at most 0.05; we use the word in this sense only. A negative genuine excess certifies *clustered structure*; the word “tree-like” is reserved for the depth test of Section 4. A positive excess means less tree-like than chance given the spectrum (Appendix A.1).

Two choices changed after pre-registration. The supremum statistic did not converge under the quadruple budget on ImageNet and disagreed with the 99.9th percentile of the same defects in both directions, so the percentile is the raw reading and the supremum is reported in the appendix. The Gaussian-coefficient null does not reproduce the sample spectrum it claims to match when n is small relative to d , and it inflates the excess of low-rank clouds; the Haar null does reproduce it, so we use the Haar null and report both. Rank-based calibration is what gives permutation scores their finite-sample validity, so a z -score is kept only as a secondary summary, and the null is aggregation-aware in the sense of Gröger et al. (2026, §5): the same percentile of the same sampled defects is evaluated on every replicate (Table 33).

Projection and tasks. A parameter-free map centers the features, scales the 95th-percentile norm to radius $1/\sqrt{2}$ and preserves directions, so Euclidean, Poincaré and cosine readouts can be compared on the intact representation; nearest-centroid and 5-way 5-shot classification score them for Section 6 (Appendix A.10). With the instrument in place, we read the premise where it is read.

4 FINDINGS I: LATENT HYPERBOLICITY, CALIBRATED

If standard models were latently hyperbolic, the instrument should find an excess on their image features, where the premise is read. Does it, and does anything survive on class centroids, where hierarchy lives?

model	family	class level (centroids)						sample level (images)	
		$\hat{\delta}_{99.9}$ IN	excess IN	$r/200$ (p)	exc. C100	exc. C10	ρ_{WN}	exc. C100	exc. DTD
ViT-T	sup.	0.088	−0.005°	169 (.159)	−0.006°	−0.021°	+0.52	−0.002°	−0.000°
ViT-S	sup.	0.063	−0.012	200 (.005)	−0.000°	−0.056	+0.50	+0.005°	−0.006
ViT-B	sup.	0.054	−0.016	200 (.005)	+0.006°	−0.050	+0.49	+0.003°	−0.008
ViT-L	sup.	0.055	−0.016	200 (.005)	−0.020	−0.060	+0.53	−0.004°	−0.007
DINO-B	SSL	0.055	−0.020	200 (.005)	−0.032	−0.061	+0.36	−0.004°	+0.001°
DINOv2-S	SSL	0.053	−0.016	200 (.005)	−0.031	−0.071	+0.22	−0.005	+0.006°
DINOv2-B	SSL	0.034	−0.017	200 (.005)	−0.031	−0.112	+0.18	−0.009	−0.002°
DINOv2-L	SSL	0.030	−0.014	200 (.005)	−0.035	−0.128	+0.18	−0.013	−0.006
DINOv2-G	SSL	0.033	−0.008	200 (.005)	−0.040	−0.130	+0.20	−0.023	−0.008
CLIP-B	contr.	0.084	−0.007°	172 (.144)	−0.019	−0.036	+0.57	−0.006°	+0.002°
CLIP-L	contr.	0.076	−0.009	197 (.020)	−0.018	−0.043	+0.59	+0.004°	−0.001°
SigLIP-B	contr.	0.075	−0.008°	189 (.060)	−0.022	−0.050	+0.58	−0.007	−0.004°

Table 1: Clustered structure is the rule at the class level; the sample-level reading sits within null noise in most cells. One row per backbone. $\hat{\delta}_{99.9}$: raw ImageNet reading; excess: $\hat{\delta}_{99.9}$ minus the mean of 200 Haar spectrum-matched null replicates; $r/200$: replicates above the real value, with the left-tail $p = (1 + \#\{\text{null} \leq \text{real}\})/201$; genuine: Benjamini–Hochberg-corrected $p \leq 0.05$ over the 72 cells; °: not genuine; bold: less clustered than the null. ρ_{WN} : Spearman correlation of inter-centroid and WordNet distances. Sample level: the same reading on ≈ 1000 stratified training images per cell, BH over those 24 cells: genuine in 10 of 24 (full table in Appendix A.18). With centroid-resampling and estimator noise folded into the null, or the census repeated on 30 resampled centroid sets, the genuine count is 42–47 of 72 (Table 37). Cross-dataset magnitudes are not comparable. The other datasets, the supremum reading and the task columns are in Table 5; all four null $\times$ statistic verdicts per cell in Table 33.

The premise does not survive calibration. We read it first on per-image features, roughly a thousand stratified training images per cell for CIFAR-100 and DTD and all twelve backbones, where the raw reading lies in the low band the literature reports. Calibrated, it is not genuine in 14 of 24 cells, and where it is genuine the excess is small and concentrated in the self-supervised family, and that family’s structure is the angular tree of Section 5, which cosine collects. A low raw reading on image features is therefore what a random cloud of the same dimension and spectrum produces, up to a small margin. This is the reading that motivates hyperbolic methods, and it does not survive calibration as stated (Table 1, last two columns). Class centroids, where hierarchy lives, read differently, and we turn to them next.

Clustered structure is the rule. Figure 3 and Table 1 give the census of record. Almost every cell sits below its matched null, and in 49 of the 72 the gap is one that chance does not produce. On ImageNet and CIFAR-100, the two sets with a real hierarchy, the count is 18 of 24. The structure belongs to no single family or architecture: every family contributes genuine cells, two self-supervised ResNets read the same as the ViTs, and resampling the images behind each centroid does not move the reading. What does move it is the label set: the smallest backbone is genuine on one dataset only, the flat datasets sit at the resolution limit of ten classes, and a class-count control shows that the reading follows the hierarchy of the labels rather than their number (Appendix A.8). A star, however, would pass this test too, which is the subject of the next paragraph.

The count survives resampling. The census null accounts for replicate variability but not for image resampling or the estimator’s seeds. Folding all three into one spread, or repeating the census on thirty resampled centroid sets, leaves between 42 and 47 of the 72 cells genuine. On ImageNet and CIFAR-100 the count stays between 17 and 18 of the 24, and the cells that move are on the transfer sets (Appendix A.15).

A star passes the census. The excess measures structure beyond second moments, and a pure star already has such structure. A synthetic star of thirty clusters with no hierarchy scores an excess as deep as any real cell. We call this the *star caveat*. The excess therefore measures how classes are arranged, not the geometry they are embedded in, and a negative excess says clustered, not deep. Whether the arrangement is deeper than a star is a separate question, which needs a test of its own (Appendix A.5). That test is the subject of the next paragraph.




Figure 4: **The depth test certifies hierarchy above the superclasses in 4 of 12 ImageNet backbones, raises no false alarm on real clouds, and misses implanted depth at ImageNet’s noise level.** Plotted: (a) the excess under the hub-randomizing null for the real ImageNet centroids and for their matched stars, isotropic as hollow circles and anisotropic as squares, with the z of real minus anisotropic star under each backbone, red at $z \leq -2$; (b) the z of the same test on the real clouds with the hub arrangement replaced by an implanted two-level tree of strength s , dashed for the same clouds with their within-cluster spread shrunk into the validated range; the synthetic sweep is in Appendix A.5.

Depth is certified on ImageNet and not on CIFAR-100. The depth test compares the real centroids with a matched star: Gaussian hubs at the real hub radius and, within each cluster, a Haar sample with the cluster’s own covariance, both read under a null that randomizes the hubs and keeps every cluster intact. Because the null keeps the clusters, the test certifies hierarchy above the labelled clusters only. On ImageNet the real cloud is more hierarchical than its matched star in 4 of 12 backbones, three supervised ViTs and DINOv2-L, and never less. The other eight are not detected, which is not the same as not hierarchical. CIFAR-100, with clusters of a few points, lies outside the regime the test covers and its readings stay unvalidated. An earlier isotropic star produced false alarms and is withdrawn (Figure 4a). How much depth the test can see on real clouds is the subject of the next paragraph.

On real clouds the test is conservative and weak. We measured the test on the real ImageNet clouds themselves, replacing the hub arrangement by an implanted two-level tree of graded strength while keeping every within-cluster offset, and the census cannot tell the implanted tree from the star. The test raises no false alarm, declaring none of the sixty zero-strength runs hierarchical. A clean implanted tree is detected in none of the twelve backbones, and the misses are a matter of noise level: the within-cluster spread of the real clouds is 1.3 to 3.9 times their between-hub spread, the synthetic sweep covered ratios below one, and once the real offsets are shrunk into that range every backbone is detected. The test is therefore conservative and weak at ImageNet’s noise level: its positive verdicts stand, and a backbone it does not certify is not thereby flat. A trained control, a ViT-B fine-tuned with a hierarchical objective, is inconclusive, as Appendix A.5 reports (Figure 4b). Whether imposing the geometry creates detectable depth is the next question.

Imposing the geometry does not create depth. MERU embeds images and text on the Lorentz hyperboloid with an entailment objective, and its released Euclidean twin is a CLIP baseline with the same backbone, data and recipe minus the hyperbolic lift, so the pair isolates the geometry (Desai et al., 2023). Passed through the instrument unchanged, the two read the same excess on ImageNet images, and MERU’s reading in its own Lorentz metric equals its Euclidean reading to the fourth decimal. Under the depth test, conservative and weak at this noise level, neither model is detected as more hierarchical than a matched star, and MERU is never ahead of its twin. Training in hyperbolic space therefore neither creates the clustering nor adds detectable depth (Table 11). If the geometry does not explain the clustering, the training objective might, and neural collapse is the first candidate.

Neural collapse is the flat limit, not the reading. Neural collapse predicts that the class means of a supervised classifier converge to a simplex equiangular tight frame (Pappas et al., 2020). An exact frame is the extreme case of the star caveat, δ zero with no hierarchy at all, and our reading differs from it in two ways. The arrangement we certify is not flat, because class means cluster around superclass hubs that a frame would erase, and cutting the centroid dendrograms recovers the CIFAR-100 superclasses. The same clustered structure also appears in models trained without class

labels, on transfer datasets and in causal language models, outside the setting where neural collapse is stated, and the depth test asks what neural collapse does not, whether there is hierarchy above the hubs, answering yes for four ImageNet backbones (Appendix A.16). What moves the excess, if not the objective’s flat limit, is scale and training, which we take up last.

Scale deepens the excess for one family only. Figure 3b shows that scale does not deepen the excess in general, with one exception: DINOv2, whose excess deepens with size on CIFAR-10 and, weakly, on CIFAR-100. Training moves the raw reading within an architecture: fine-tuning on a task with no class hierarchy raises it, it falls across transformer depth, and shuffling the embedding-to-label correspondence leaves it unchanged (Appendix A.6). Together, these results say that what survives at the class level is clustered structure that is genuine in most cells, rarely deep, and not created by the geometry of the embedding space. Whether that structure is shared across models, and whose tree it is, is the subject of the next section.

5 FINDINGS II: WHOSE TREE

If the measured trees were a shared representation of reality, models should agree with each other and with human taxonomy. Whether they appear to depends on how the trees are measured.

The naive map manufactures an island. We extract each model’s hierarchy as a dendrogram over the shared centroids and compare hierarchies by the adjusted Rand index between cut partitions and by the cophenetic correlation between tree and distances. Under the naive configuration, Euclidean distances with average linkage, the map is spectacular: supervised ViTs and contrastive VLMs form one block, and the label-free DINOv2 family appears as an island that agrees with no supervised model. We call this the *island*. Its cuts are degenerate: DINOv2’s Euclidean dendrograms chain, almost every class falls into one cluster, and the index is forced to zero by construction. A configuration must therefore be selected by a criterion fixed without reference to the agreements being compared: we require the highest cophenetic fidelity among configurations whose reference cuts are non-degenerate for every model, which selects cosine-average on ImageNet and cosine-complete on CIFAR-100. The criterion is applied per dataset, because degeneracy is a property of the dataset, and it never sees the agreements it selects for. Under every non-degenerate configuration the gap stays moderate and never returns to the island, and the configuration it selects on DBpedia is the one that degenerates on ImageNet, which is why it cannot be transferred. Figure 5 shows the map before and after this control (Appendix A.9). Under the selected configurations the island has to be re-read.

The island is an artifact, and sharing is graded. Under the selected configurations three measures agree. At the cut, the island shrinks from categorical to moderate: DINOv2 agrees with the block at 0.38 against 0.48 within the block. In cophenetic correlation a gap remains. In triplet agreement, which needs no cut, the gap closes entirely under angular configurations and persists under Euclidean ones. Across every non-degenerate configuration, label-free models recover the human taxonomy partially and supervision deepens the resemblance: paired within each admissible configuration, the label-supervised ViT recovers more than each DINOv2 model in 11 of 12 comparisons (Appendix A.9). The gap is a property of the measure as much as of the models, and its persistence under Euclidean measures alone points to where the self-supervised structure lives.

The self-supervised tree lives in angles. DINOv2’s semantics concentrate in the angular component of its representation: for DINOv2-L, sibling triplets are resolved far better under cosine than under Euclidean distance. We call this structure the *angular tree*. It is why DINOv2’s Euclidean dendrograms chain into the island and why the triplet gap closes only under angular configurations. The census agrees when we repeat it on cosine geometry, with L2-normalized centroids, geodesic distances and the same Haar null on the normalized cloud: the verdict matches the Euclidean census in 65 of 72 cells, and DINOv2 on ImageNet is genuine under both readings. We report the cosine census as robustness, never as the record (Table 35). Where the structure lives decides which readout collects it, the concern of Section 6. First we ask whose taxonomy the trees recover.

WordNet alignment follows supervision, not tree-likeness. The correlation between inter-centroid distances and WordNet tree distances on ImageNet orders the families by supervision: highest for the contrastive VLMs, then the supervised ViTs, and lowest for DINOv2, with shuffle and




Figure 5: **The island vanishes once the clustering step is controlled.** Plotted: pairwise ARI between the dendrogram cuts of the 12 vision models on ImageNet under the naive configuration (a) and the selected one (b), with the mean DINOv2-vs-block ARI above each panel.

Gaussian controls at zero. DINOv2 is nevertheless genuine on ImageNet and both CIFAR sets, and cutting centroid dendrograms recovers the true CIFAR-100 superclasses at an adjusted Rand index of up to 0.61, against zero for permuted labels. Across the three families, WordNet content and robustness of the excess anti-correlate: the supervised and contrastive families are the most aligned and the most fragile, DINOv2 the least aligned and the most robust (Table 1). Alignment with the human taxonomy is therefore something training puts in, not something tree-likeness reveals, which raises the question of circularity.

The recovery is not WordNet circularity. ImageNet classes are WordNet synsets, so recovering WordNet there could be circular. Three controls bound the circularity. WordNet-coherent groupings read a lower raw δ than random ones in 11 of 12 models. Single-image distances still correlate with WordNet, up to a correlation of +0.26. And the decoupling of alignment from robustness is the opposite of what label injection would produce. Repeating the tree map on DBpedia Classes, a three-level ontology independent of WordNet, closes the door: no configuration degenerates, the three text embedders recover the true level-2 partition in every configuration, and they agree with each other (Table 42). The recovered hierarchy is real, shared and configuration-robust, and the fragility above belongs to the family and the concept set together, which text makes explicit.

In text, clustered structure depends on recipe, scale and probe. On DBpedia, a real three-level text hierarchy, the sentence embedders read a genuine excess with every replicate above the real value. On the harder probe, the thousand ImageNet class names, the same embedders are not genuine, despite raw readings indistinguishable from vision backbones whose excess is. Within causal language models on the class-name probe, the two largest GPT-2 models are genuine under every construction of the null and statistic and under cosine, whereas the two smallest flip across readings: robust at the larger scales, fragile at the smaller. Pythia and OLMo-1B are genuine at every scale tested, and extraction protocol and prompt template move the fragile values but not the scale ordering (Appendix A.7). The trees, then, are shared in content only partially, and whether they are shared in local structure at all is the last question of this section.

Local agreement survives calibration. We apply the permutation calibration of Gröger et al. (2026) to the 66 model pairs: permutations of the class correspondence give the null distribution, its upper quantile gives the threshold, and the calibrated score is the excess over that threshold rescaled to keep the maximum at one. Mutual-kNN agreement survives calibration with every pair significant, and so does centroid CKA, while the Poincaré reading improves neither: centroids distill a relational base that models share locally, in who is near whom, even where their trees differ (Table 24). Together, these findings give the thesis its short form. Clustered, occasionally hierarchical, moderately shared: not one common tree, and no license for curvature.

6 CONSEQUENCES FOR IMPOSING CURVATURE

If the raw reading cannot be read on its own, what should a practitioner read before imposing curvature? And if the structure that survives is clustered and angular, which readout collects it?

The raw reading cannot select a curvature. Khrulkov et al. (2020) set the curvature of the Poincaré ball from the measured relative hyperbolicity as $c = (0.144/\delta_{\text{rel}})^2$ with $\delta_{\text{rel}} = 2\delta/\text{diam}$. Since the raw reading is dominated by the dimension and spectrum confounds, such a curvature tracks the backbone’s width rather than its hierarchy. Structureless Gaussian clouds would receive a curvature from 0.48 at $d=192$ to 2.5 at $d=1536$. The excess is the quantity to read instead (Table 1).

The excess is small in absolute terms and large against the null. On ImageNet centroids it runs from -0.005 to -0.020 . That is 1 to 36 times the null’s standard deviation at that class count. On image features it sits within noise in most cells. What to read instead is the next paragraph.

Measure before imposing. For a frozen backbone the instrument answers three questions in minutes. Is there structure beyond a spectrum-matched null? The excess and its rank say. Is it deeper than a star? The depth test says, conservative and weak at ImageNet’s noise level, so that a positive verdict certifies and a negative one says nothing. Where does it live, in angles or in distances? The cosine reading of Section 5 says. In this census the answers are usually yes, rarely, and, for the self-supervised family, in angles, and a change of readout collects that structure at no cost.

A zero-cost readout collects what is there. Changing the readout metric on the intact representation costs nothing, and cosine collects the angular tree. The Poincaré readout adds from $+0.1$ to $+1.3$ pp over cosine for the contrastive VLMs on prototype tasks. For the other families its effect is inconsistent, from -1.7 to $+1.1$ pp. On flat-label datasets no metric helps, and sample-level tasks never benefit. The full study is in Appendix A.10, and the discussion states the three acts together.

7 DISCUSSION AND LIMITATIONS

What the paper establishes. Read correctly, foundation models organize classes into clustered structure that is occasionally hierarchical and moderately shared; they do not converge to one common tree, and their raw tree-likeness is not evidence for hyperbolic geometry. In three acts: the premise of latent hyperbolicity is confounded and its sample-level reading sits within null noise; what survives at the class level is clustered structure in 49 of 72 cells, hierarchy above the superclasses in 4 of 12 ImageNet backbones, no detected depth from imposing the geometry, and a moderately shared, supervision-dependent recovery of the human taxonomy, angular for the self-supervised family; and what follows is that curvature should be measured rather than assumed and a zero-cost readout collects what is there.

Open questions. Why the structure is clustered but rarely certified as deep, and why the self-supervised tree is angular, are open: every recipe recovers the superclass hubs, yet only four backbones place them detectably beyond a star.

Limitations. (i) The excess certifies structure beyond second moments, star-like clustering included; the depth test certifies hierarchy above its frame only, is conservative and weak at ImageNet’s noise level, leaves hierarchy within the clusters untested, and its one trained positive control is inconclusive (Appendix A.5). (ii) The census is class-centroid geometry on the datasets of Table 1. (iii) Causal-LM values, not their scale ordering, are template-conditioned. (iv) The objective-geometry link is correlational and rests chiefly on DINOv2’s scale range, as Appendix A.13 shows. (v) Downstream scope is prototype-based metric tasks with modest gains. (vi) The descriptor ξ tracks norm dispersion rather than curvature, as Appendix A.21 reports. (vii) A causal LM’s class representation depends on the extraction protocol as much as on the model. (viii) The sample-level reading covers two datasets and one image budget.

ETHICS STATEMENT

This work analyzes publicly released models and standard public datasets, involves no human subjects and no personal data, and proposes an analysis methodology and a metric-selection heuristic. We do not foresee ethical risks beyond those already associated with the underlying models and datasets.

REPRODUCIBILITY STATEMENT

Every number in the paper is annotated in the source with the result file that produces it; all tables are generated by script from those files, and the full pipeline (feature extraction, diagnostics, nulls, calibration, tasks, statistics, figures) with fixed seeds will be released. Estimator settings are stated in §3. The instrument ships as one script (ANONYMIZED-REPO) that reproduces any cell of Table 1 from its centroid matrix; its output on two cells is checked against the table at every freeze. Computing δ_{norm} on 1000 centroids takes seconds on CPU; the full task grid runs in under 30 minutes on one consumer GPU.

THE USE OF LARGE LANGUAGE MODELS

LLM-based assistants were used, under continuous author direction and review, to draft and edit text, to write analysis scripts, and to run the experiment pipeline on the authors’ hardware. All experimental designs, results, verifications, and final claims were checked and approved by the authors, who take full responsibility for the content of this paper.

REFERENCES

- Aaron B. Adcock, Blair D. Sullivan, and Michael W. Mahoney. Tree-like structure in large social and information networks. In *IEEE International Conference on Data Mining (ICDM)*, pp. 1–10, 2013.
- Morris Alper and Hadar Averbuch-Elor. Emergent visual-semantic hierarchies in image-text representations. In *European Conference on Computer Vision (ECCV)*, 2024.
- Alessio Ansuini, Alessandro Laio, Jakob H. Macke, and Davide Zoccolan. Intrinsic dimension of data representations in deep neural networks. In Hanna M. Wallach, Hugo Larochelle, Alina Beygelzimer, Florence d’Alché-Buc, Emily B. Fox, and Roman Garnett (eds.), *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*, pp. 6109–6119, 2019. URL <https://proceedings.neurips.cc/paper/2019/hash/cfcce0621b49c983991ead4c3d4d3b6b-Abstract.html>.
- Mina Ghadimi Atigh, Julian Schoep, Erman Acar, Nanne van Noord, and Pascal Mettes. Hyperbolic image segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 4453–4462, June 2022.
- Ahmad Bdeir, Kristian Schwethelm, and Niels Landwehr. Fully hyperbolic convolutional neural networks for computer vision. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net, 2024. URL <https://openreview.net/forum?id=ekz1hN5QNh>.
- Ines Chami, Zhitao Ying, Christopher Ré, and Jure Leskovec. Hyperbolic graph convolutional neural networks. In Hanna M. Wallach, Hugo Larochelle, Alina Beygelzimer, Florence d’Alché-Buc, Emily B. Fox, and Roman Garnett (eds.), *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*, pp. 4869–4880, 2019. URL <https://proceedings.neurips.cc/paper/2019/hash/0415740eaa4d9decbc8da001d3fd805f-Abstract.html>.
- Weize Chen, Xu Han, Yankai Lin, Hexu Zhao, Zhiyuan Liu, Peng Li, Maosong Sun, and Jie Zhou. Fully hyperbolic neural networks. In Smaranda Muresan, Preslav Nakov, and Aline

- Villavicencio (eds.), *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 5672–5686, Dublin, Ireland, May 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.acl-long.389. URL <https://aclanthology.org/2022.acl-long.389/>.
- Karan Desai, Maximilian Nickel, Tanmay Rajpurohit, Justin Johnson, and Ramakrishna Vedantam. Hyperbolic image-text representations. In *Proceedings of the 40th International Conference on Machine Learning, ICML’23*. JMLR.org, 2023.
- Aleksandr Ermolov, Leyla Mirvakhabova, Valentin Khrulkov, Nicu Sebe, and Ivan V. Oseledets. Hyperbolic vision transformers: Combining improvements in metric learning. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2022, New Orleans, LA, USA, June 18-24, 2022*, pp. 7399–7409. IEEE, 2022. doi: 10.1109/CVPR52688.2022.00726. URL <https://doi.org/10.1109/CVPR52688.2022.00726>.
- Hervé Fournier, Anas Ismail, and Antoine Vigneron. Computing the gromov hyperbolicity of a discrete metric space. *Information Processing Letters*, 115(6-8), 2015.
- Octavian Ganea, Gary Becigneul, and Thomas Hofmann. Hyperbolic neural networks. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett (eds.), *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc., 2018. URL https://proceedings.neurips.cc/paper_files/paper/2018/file/dbab2adc8f9d078009ee3fa810bea142-Paper.pdf.
- Fabian Gröger, Song Wen, and Maria Brbić. Revisiting the platonic representation hypothesis: An aristotelian view. In *International Conference on Machine Learning (ICML)*, 2026. arXiv:2602.14486.
- M. Gromov. *Hyperbolic Groups*, pp. 75–263. Springer New York, New York, NY, 1987. ISBN 978-1-4613-9586-7. doi: 10.1007/978-1-4613-9586-7_3. URL https://doi.org/10.1007/978-1-4613-9586-7_3.
- Neil He, Rishabh Anand, Hiren Madhu, Ali Maatouk, Smita Krishnaswamy, Leandros Tassioulas, Menglin Yang, and Rex Ying. HELM: Hyperbolic large language models via mixture-of-curvature experts. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025. URL <https://openreview.net/forum?id=RnbJPkakkM>.
- Minyoung Huh, Brian Cheung, Tongzhou Wang, and Phillip Isola. Position: The platonic representation hypothesis. In Ruslan Salakhutdinov, Zico Kolter, Katherine A. Heller, Adrian Weller, Nuria Oliver, Jonathan Scarlett, and Felix Berkenkamp (eds.), *Forty-first International Conference on Machine Learning, ICML 2024, Vienna, Austria, July 21-27, 2024*, Proceedings of Machine Learning Research, pp. 20617–20642. PMLR / OpenReview.net, 2024. URL <https://proceedings.mlr.press/v235/huh24a.html>.
- W. Sean Kennedy, Onuttom Narayan, and Iraj Saniee. On the hyperbolicity of large-scale networks. *arXiv preprint arXiv:1307.0031*, 2013.
- Valentin Khrulkov, Leyla Mirvakhabova, Evgeniya Ustinova, Ivan Oseledets, and Victor Lempitsky. Hyperbolic image embeddings. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2020.
- A. Sophia Koepke, Daniil Zverev, Shiry Ginosar, and Alexei A. Efros. Back into plato’s cave: Examining cross-modal representational convergence at scale, 2026. arXiv preprint.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey E. Hinton. Similarity of neural network representations revisited. In Kamalika Chaudhuri and Ruslan Salakhutdinov (eds.), *Proceedings of the 36th International Conference on Machine Learning, ICML 2019, 9-15 June 2019, Long Beach, California, USA*, Proceedings of Machine Learning Research, pp. 3519–3529. PMLR, 2019. URL <http://proceedings.mlr.press/v97/kornblith19a.html>.
- Onuttom Narayan and Iraj Saniee. Large-scale curvature of networks. *Physical Review E*, 84(6): 066108, 2011.

- Bogdan Nica and Ján Špakula. Strong hyperbolicity. *Groups, Geometry, and Dynamics*, 10(3), 2016. arXiv:1408.0250.
- Maximilian Nickel and Douwe Kiela. Poincaré embeddings for learning hierarchical representations. In Isabelle Guyon, Ulrike von Luxburg, Samy Bengio, Hanna M. Wallach, Rob Fergus, S. V. N. Vishwanathan, and Roman Garnett (eds.), *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA*, pp. 6338–6347, 2017. URL <https://proceedings.neurips.cc/paper/2017/hash/59dfa2df42d9e3d41f5b02bfc32229dd-Abstract.html>.
- Yann Ollivier. Ricci curvature of Markov chains on metric spaces. *Journal of Functional Analysis*, 256(3):810–864, 2009.
- Avik Pal, Max van Spengler, Guido Maria D’Amely di Melendugno, Alessandro Flaborea, Fabio Galasso, and Pascal Mettes. Compositional entailment learning for hyperbolic vision-language models. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=3i13Gev2hV>.
- Vardan Papyan, X. Y. Han, and David L. Donoho. Prevalence of neural collapse during the terminal phase of deep learning training. *Proceedings of the National Academy of Sciences*, 117(40):24652–24663, 2020. doi: 10.1073/pnas.2015509117. URL <https://www.pnas.org/doi/abs/10.1073/pnas.2015509117>.
- Kiho Park, Yo Joong Choe, Yibo Jiang, and Victor Veitch. The geometry of categorical and hierarchical concepts in large language models. *arXiv preprint arXiv:2406.01506*, 2024.
- Phillip Pope, Chen Zhu, Ahmed Abdelkader, Micah Goldblum, and Tom Goldstein. The intrinsic dimension of images and its impact on learning. In *9th International Conference on Learning Representations, ICLR 2021, Virtual Event, Austria, May 3-7, 2021*. OpenReview.net, 2021. URL <https://openreview.net/forum?id=XJk19XzGq2J>.
- Rik Sarkar. Low distortion delaunay embedding of trees in hyperbolic plane. In Marc J. van Kreveld and Bettina Speckmann (eds.), *Graph Drawing - 19th International Symposium, GD 2011, Eindhoven, The Netherlands, September 21-23, 2011, Revised Selected Papers*, Lecture Notes in Computer Science, pp. 355–366. Springer, 2011. doi: 10.1007/978-3-642-25878-7_34. URL https://doi.org/10.1007/978-3-642-25878-7_34.
- Aditya Sinha, Siqi Zeng, Makoto Yamada, and Han Zhao. Learning structured representations with hyperbolic embeddings. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang (eds.), *Advances in Neural Information Processing Systems*, volume 37, pp. 91220–91259. Curran Associates, Inc., 2024. doi: 10.52202/079017-2895. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/a5d2da376bab7624b3caeb9f78fcaa2f-Paper-Conference.pdf.
- Alexandru Tifrea, Gary Bécigneul, and Octavian-Eugen Ganea. Poincaré glove: Hyperbolic word embeddings. In *International Conference on Learning Representations (ICLR)*, 2019.
- Menglin Yang, Ram Samarth B B, Aosong Feng, Bo Xiong, Jiahong Liu, Irwin King, and Rex Ying. Hyperbolic fine-tuning for large language models. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025. URL <https://openreview.net/forum?id=TkEdQv0bXB>.

A FULL NUMERICAL TABLES

Every table is generated by script directly from the released per-cell result files; no value is hand-transcribed.

A.1 WHAT EACH NULL PRESERVES

null	preserves	destroys
Gaussian, spectrum-matched	n, d , mean, covariance spectrum	all higher-order structure
PC-permutation	each principal component’s marginal	dependence across components
Cluster-preserving (tested, unused)	each cluster’s geometry, hub assignment	hub arrangement (invalid depth test: App. A.5)
Label shuffle (ablation)	the point cloud itself	embedding–label correspondence

Table 2: The nulls of the instrument and what each one preserves and destroys; “genuine” in the text always means beyond the first of these. In plain terms, the spectrum null is a cloud with the same shape as the real one and no structure inside it; the hub null of Section 4 keeps every cluster of the real cloud and only rearranges where the clusters sit.

A.2 ESTIMATOR CALIBRATION

space	absolute δ	δ_{norm}
balanced binary tree (depth 10)	0.000	0.000
$\mathbb{H}^2$ ($K = -1$), region radius $R=2/4/8/16$	0.65/0.69/0.69/0.69	0.162/0.087/0.043/0.022
uniform S^{99} (chord / geodesic)	0.25/0.37	0.143/0.179
iid Gaussian ($n=1000$), $d=192, \dots, 1536$	—	0.104/0.081/0.074/0.061/0.055/0.046

Table 3: **A low raw δ is not evidence of hierarchy.** Calibration on reference geometries: the estimator recovers the tree zero and the four-point constant of $\mathbb{H}^2$, spheres stay high, and iid Gaussians drift low as d grows, the dimension confound. Gaussian row from the released control file; tree/ $\mathbb{H}^2$ /sphere rows re-verified synthetically at every freeze.

A.3 MODEL PANEL

Class centroids average 100 cached training images per class on ImageNet and the full cached split, 80–6,000 images per class, on the other datasets. Class centroids average 100 cached training images per class on ImageNet and the full cached split, 80–6,000 images per class, on the other datasets. Class centroids average 100 cached training images per class on ImageNet and the full cached split, 80–6,000 images per class, on the other datasets. Feature extraction: vision backbones use the encoder’s default pooled embedding (timm, num_classes=0, i.e. CLS or global pool per architecture); causal LMs use the last-token hidden state, extracted one prompt at a time (padding-free: left-padded batching alters GPT-2’s hidden states, Table 15); text embedders use their recommended pooling (CLS for BGE, mean for GTE/E5, last token for GTE-Qwen2); CLIP/SigLIP text towers use their text encoders.

type	model	params
Supervised ViT (vision)	ViT-T (i21k)	5M
Supervised ViT (vision)	ViT-S (i21k)	22M
Supervised ViT (vision)	ViT-B (i21k)	86M
Supervised ViT (vision)	ViT-L (i21k)	307M
SSL (vision)	DINO-B	86M
SSL (vision)	DINOv2-S	22M
SSL (vision)	DINOv2-B	86M
SSL (vision)	DINOv2-L	307M
SSL (vision)	DINOv2-G	1.1B
Contrastive (vision)	CLIP-B	86M
Contrastive (vision)	CLIP-L	307M
Contrastive (vision)	SigLIP-B	86M
Causal LM	GPT-2 S	117M
Causal LM	GPT-2 M	345M
Causal LM	GPT-2 L	774M
Causal LM	GPT-2 XL	1.5B
Causal LM	Pythia-410M	410M
Causal LM	Pythia-1B	1.0B
Causal LM	Pythia-2.8B	2.8B
Causal LM	OLMo-1B	1.0B
Causal LM	OLMo-7B	7.0B
Text embedder	BGE-base	110M
Text embedder	BGE-large	335M
Text embedder	GTE-base	110M
Text embedder	GTE-large	335M
Text embedder	GTE-Qwen2-1.5B	1.5B
Text embedder	E5-base	110M
Text embedder	E5-large	335M

Table 4: The full model panel: 12 vision backbones (all in the census; the 10 with cached per-dataset features, i.e. all but DINO-B and SigLIP-B, form the task grid), 9 causal LMs, and 7 text embedders.

A.4 VISION CENSUS: EXCESS AND SIGNIFICANCE PER CELL

model	sup. IN (r)	exc. DTD	exc. FMNIST	exc. MNIST	best-R (pp)	metric
ViT-T	−0.020 (199)	−0.014	+0.018°	−0.026°	+0.58	either
ViT-S	−0.020 (200)	−0.019	−0.033	−0.037°	+0.49	either
ViT-B	−0.022 (200)	−0.021	−0.042	−0.013°	+0.51	either
ViT-L	−0.027 (200)	−0.031	−0.051	−0.040°	+0.87	either
DINO-B	+0.000 (96)	−0.027	−0.007°	−0.034°	—	—
DINOv2-S	+0.017 (0)	−0.018	−0.031°	−0.004°	+0.68	H
DINOv2-B	+0.018 (0)	−0.035	−0.062	−0.017°	+1.17	cos
DINOv2-L	+0.006 (0)	−0.042	−0.035	−0.014°	+1.84	cos
DINOv2-G	+0.020 (0)	−0.039	−0.046	−0.022°	+2.54	cos
CLIP-B	−0.023 (199)	−0.024	−0.035°	−0.001°	+0.67	H
CLIP-L	−0.016 (198)	−0.027	−0.037°	−0.008°	+0.65	H
SigLIP-B	−0.015 (194)	−0.023	−0.042	−0.044	—	—

Table 5: Census columns moved out of Table 1. sup. IN: ImageNet excess of the supremum statistic under the same Haar null (200 replicates; r = replicates above the real value; Table 33); exc. DTD/FMNIST/MNIST: per-dataset record excess on the remaining sets; best-R: few-shot advantage of the best zero-cost metric over Euclidean, mean over the four hierarchical datasets; metric: which of Poincaré (H) or cosine collects it (either: within 0.15pp; —: outside the 10-model task grid).

model	d	$\hat{\delta}$	excess (z)	reps above
Barlow Twins (ResNet-50)	2048	0.075	−0.039 (−4.8)	20/20
BYOL (ResNet-50)	2048	0.073	−0.020 (−2.4)	20/20

Table 6: Backbones outside the ViT census, ImageNet centroid store: two ResNet-50 self-supervised models, under the spectrum null (20 replicates, percentile rank).

A.5 WHAT THE EXCESS CERTIFIES: SYNTHETIC CALIBRATION




Figure 6: **The depth test on CIFAR-100 ($K=20$, five points per cluster): the unvalidated regime.** Excess B under the hub-randomizing null for the real centroids (filled) and their matched stars (hollow circle: isotropic, withdrawn; square: anisotropic); the number under each backbone is the z of real minus anisotropic star. In this regime pure synthetic stars are declared more hierarchical in 17% of runs (Table 9), so these readings are reported without certification.

Real-data readings by regime. ImageNet ($n=1000$, WordNet $K=30$) lies in the validated regime: the real cloud is more hierarchical than its anisotropic matched star beyond the noise ($z \leq -2$) in 4 of 12 backbones (ViT-S, ViT-B, ViT-L, DINOv2-L), never less (Table 8, Figure 4a). CIFAR-100 ($n=100$, $K=20$) lies in the unvalidated regime, where pure stars are declared more hierarchical in 17% of runs; its readings (3 of 12 with $z \leq -2$, none with $z \geq +2$) are reported without certification (Figure 6). Validity domain of the test as run: $n \approx 1000$; at $n=100$ the false-alarm rate is 0% with $K=6$ (17 points per cluster) and 17% with $K \geq 12$, so the test is not validated below about ten points per cluster. The bar was fixed before the sweep; the restriction to the validated regime was adopted after it.

synthetic cloud ($n=1000$, $d=768$)	$\hat{\delta}$	spectrum null A		hub-randomizing null B	
		Gaussian	Haar	Gaussian	Haar
30-cluster star, tight (0.1)	0.055	-0.115	-0.115	-0.126	-0.058
30-cluster star, mid (0.3)	0.059	-0.104	-0.103	-0.114	-0.049
30-cluster star, loose (0.6)	0.061	-0.084	-0.082	-0.090	-0.033
2-level hierarchy 6×5 (0.3)	0.044	-0.143	-0.143	-0.113	-0.112

Table 7: Calibration of the nulls on synthetic clouds (excess, means over seeds; within/between-cluster noise ratio in parentheses). A pure star (Gaussian clusters around Gaussian centers, no hierarchy) sits far below the spectrum null A under either construction (Gaussian coefficients, the paper’s; Haar-rotated coefficients, exact sample spectrum): null A certifies clustering, not depth. The hub-randomizing null B, which keeps every cluster intact and resamples only the hub configuration, also reports the star as “hierarchical”: a near-regular simplex of hubs already minimizes $\hat{\delta}$, so any hub randomization raises it. Null B is therefore not a valid depth test as constructed and is not used in the paper; a star-calibrated version (excess B relative to a matched star with the same cluster count and tightness) may still serve as a depth test, since under the Haar construction the two-level hierarchy sits 2–3 $\times$ below the star. That version is run on the real backbones in Table 8. Under the plain Gaussian null the mid-radius star scores -0.104 and the 6×5 hierarchy -0.14: a star and a two-level tree differ by a fraction of either’s excess, which is why depth cannot be read from the spectrum excess alone.

model	CIFAR-100 ($K=20$)		ImageNet ($K=30$)	
	isotropic star	anisotropic star	isotropic star	anisotropic star
ViT-T	+0.000 (+0.0)	-0.025 (-1.4)	-0.015 (-1.5)	-0.018 (-1.5)
ViT-S	+0.036 (+2.3)	-0.008 (-0.5)	-0.021 (-1.9)	-0.028 (-2.4)
ViT-B	+0.063 (+4.5)	+0.003 (+0.2)	-0.026 (-2.9)	-0.033 (-3.6)
ViT-L	+0.040 (+4.2)	-0.017 (-1.5)	-0.024 (-3.2)	-0.033 (-4.0)
DINO-B	+0.008 (+0.9)	-0.038 (-4.2)	-0.002 (-0.2)	-0.009 (-0.9)
DINOv2-S	+0.006 (+0.6)	-0.020 (-1.6)	+0.011 (+0.9)	+0.008 (+0.7)
DINOv2-B	+0.041 (+5.4)	+0.010 (+0.9)	-0.011 (-0.9)	-0.016 (-1.2)
DINOv2-L	+0.042 (+6.8)	+0.013 (+1.2)	-0.024 (-2.3)	-0.028 (-2.4)
DINOv2-G	+0.026 (+3.8)	-0.007 (-0.7)	-0.024 (-1.6)	-0.029 (-1.8)
CLIP-B	+0.019 (+1.8)	-0.024 (-2.0)	-0.006 (-0.7)	-0.016 (-2.0)
CLIP-L	+0.036 (+5.7)	-0.013 (-2.3)	+0.004 (+0.5)	-0.008 (-0.8)
SigLIP-B	+0.027 (+2.8)	-0.021 (-2.2)	+0.001 (+0.0)	-0.011 (-0.9)

Table 8: The matched-star depth test on the real backbones: the star’s shape decides the verdict. The test is validated at $n=1000$ only (Table 9); CIFAR-100 readings ($n=100$) are reported without certification. Certified on ImageNet ($n=1000$, the validated regime): 4 of 12 backbones (ViT-S, ViT-B, ViT-L, DINOv2-L) with z from -2.4 to -4.0 ; the withdrawn isotropic star reached $z = +6.8$ on CIFAR-100. Depth = excess B of the real centroids minus that of a matched star (negative = more hierarchical above the frame than the star), with z against the combined spread; 10 star seeds; hub-randomizing Haar null. Isotropic star: Gaussian clouds with each superclass’s RMS spread (the earlier, withdrawn construction). Anisotropic star: within each superclass a Haar sample with the cloud’s own covariance. Frames: CIFAR-100 coarse labels; ImageNet WordNet cut. K sweep: CIFAR-100 $K=5$: iso 0/12 at $z \leq -2$, 7/12 at $z \geq +2$; aniso 0/12 at $z \leq -2$, 0/12 at $z \geq +2$. CIFAR-100 $K=10$: iso 0/12 at $z \leq -2$, 9/12 at $z \geq +2$; aniso 1/12 at $z \leq -2$, 0/12 at $z \geq +2$. CIFAR-100 $K=20$: iso 0/12 at $z \leq -2$, 8/12 at $z \geq +2$; aniso 3/12 at $z \leq -2$, 0/12 at $z \geq +2$. ImageNet $K=10$: iso 0/12 at $z \leq -2$, 0/12 at $z \geq +2$; aniso 2/12 at $z \leq -2$, 0/12 at $z \geq +2$. ImageNet $K=30$: iso 3/12 at $z \leq -2$, 0/12 at $z \geq +2$; aniso 4/12 at $z \leq -2$, 0/12 at $z \geq +2$. ImageNet $K=60$: iso 3/12 at $z \leq -2$, 0/12 at $z \geq +2$; aniso 4/12 at $z \leq -2$, 0/12 at $z \geq +2$.




Figure 7: Power of the depth test on synthetic hierarchies at ImageNet’s class count. Plotted: power ($z \leq -2$) on synthetic two- and three-level hierarchies with the frame at the leaf clusters (solid, one line per K) or at the top level (dotted), against the within/between noise ratio, with the false-alarm rates of pure stars dashed.

frame	n	K	power @ ratio 0.1	0.3	0.6	star $z \leq -2$	star $z \geq +2$
top-level, isotropic star	100	6	0.00	0.00	0.00	0.00	0.17
top-level, isotropic star	100	12	0.00	0.00	0.00	0.00	0.13
top-level, isotropic star	100	20	0.00	0.00	0.00	0.00	0.17
top-level, isotropic star	100	30	0.00	0.00	0.00	0.00	0.17
top-level, isotropic star	1000	6	0.00	0.00	0.00	0.00	0.17
top-level, isotropic star	1000	12	0.00	0.00	0.00	0.00	0.17
top-level, isotropic star	1000	20	0.00	0.00	0.00	0.00	0.17
top-level, isotropic star	1000	30	0.00	0.00	0.00	0.00	0.17
leaf clusters, anisotropic star	100	6	1.00	1.00	1.00	0.00	0.00
leaf clusters, anisotropic star	100	12	1.00	1.00	1.00	0.17	0.00
leaf clusters, anisotropic star	100	20	1.00	1.00	1.00	0.17	0.00
leaf clusters, anisotropic star	1000	6	1.00	1.00	1.00	0.00	0.00
leaf clusters, anisotropic star	1000	12	1.00	1.00	1.00	0.00	0.00
leaf clusters, anisotropic star	1000	20	1.00	1.00	0.90	0.00	0.00
leaf clusters, anisotropic star	1000	30	1.00	1.00	0.70	0.00	0.00

Table 9: **Power of the matched-star depth test on synthetic hierarchies.** Two- and three-level hierarchies (pooled) and pure stars, $d=768$, within/between noise ratios 0.1/0.3/0.6, isotropic and anisotropic clusters pooled, 5 seeds. Power = fraction of hierarchy runs with $z \leq -2$; star columns = false-alarm rates in each direction. Pre-set bar: power ≥ 0.8 for two- and three-level hierarchies at noise ratios ≤ 0.3 for both n , and false alarms $\leq 5\%$ in each direction. Leaf frame: met at $n=1000$ (power 1.00 at ratios ≤ 0.3 and 0.90 at 0.6; false alarms 0% and 0%), failed at $n=100$, where $K \geq 12$ (under ten points per cluster) gives 17% false alarms at $z \leq -2$. Top-level frame: power 0.00, false alarms 16% pooled. Top-level frame: the test receives the K super-cluster labels, as the earlier isotropic test (withdrawn) did on the real backbones; its hub null keeps every frame cluster intact and only rearranges the K hubs, so hierarchy below the frame is invisible by construction. Leaf frame: the test receives the finest cluster labels with the anisotropic matched star of Table 8 (configurations with more leaves than points are infeasible and omitted).

model	WordNet-30 frame of record							balanced frame (pre-registered)		
	real z	within/between	hits $s=0/0.5/1$	s^*	mean z at $s=1$	tight z at $s=0/0.5/1$		real z	hits $s=0/0.5/1$	s^*
ViT-T	-1.5	1.5	0/0/0	-	+0.1	-0.6/-1.9/-9.2		-2.1	0/0/1	-
ViT-S	-2.4	1.9	0/0/0	-	-1.2	-0.2/-3.4/-10.1		-1.7	0/0/0	-
ViT-B	-3.6	2.0	0/0/0	-	+0.4	-1.0/-3.9/-11.5		-2.8	0/0/0	-
ViT-L	-4.0	1.9	0/0/0	-	-1.4	-1.6/-4.4/-12.7		-5.6	0/0/0	-
DINO-B	-0.9	2.1	0/0/0	-	-0.3	-1.6/-5.1/-12.1		-1.0	0/0/0	-
DINOv2-S	+0.7	3.0	0/0/0	-	-0.8	-0.2/-3.1/-11.5		+0.1	0/0/0	-
DINOv2-B	-1.2	3.7	0/0/0	-	+0.4	-0.0/-3.1/-11.2		-0.3	0/0/0	-
DINOv2-L	-2.4	3.9	0/0/0	-	+0.2	+0.2/-2.7/-11.4		-1.7	0/0/0	-
DINOv2-G	-1.8	3.9	0/0/0	-	+1.3	+0.2/-3.9/-11.6		-0.9	0/0/0	-
CLIP-B	-2.0	1.3	0/0/0	-	-1.2	-1.7/-3.8/-11.4		-1.3	0/0/0	-
CLIP-L	-0.8	1.5	0/0/3	-	-2.2	-2.4/-5.8/-13.5		-0.7	0/0/0	-
SigLIP-B	-0.9	1.5	0/0/0	-	-1.6	-2.1/-5.0/-11.0		-1.1	0/0/0	-

Table 10: **Implanted depth on real ImageNet centroids: no false alarms, and low power at the real within/between spread.** The hub arrangement of each real cloud is replaced by an implanted two-level tree of strength s (six super-hubs of five, hub = $s \cdot \text{super-hub} + (1-s+0.4s) \cdot \text{own Gaussian draw}$, hub cloud rescaled to the real RMS radius) while every within-cluster offset is kept; the depth test of Table 8 runs unchanged, 5 implant seeds per s . hits: implant seeds with $z \leq -2$ (of 5); s^* : first s with ≥ 4 of 5. within/between: RMS of the within-cluster offsets over the point-weighted RMS of the hub displacements (the synthetic sweep covered 0.1–0.6). tight: the same implant after shrinking the offsets to within/between = 0.6 (2 seeds). Power (fraction of runs with $z \leq -2$) by s : $s=0$: 0.00, $s=0.25$: 0.00, $s=0.5$: 0.00, $s=0.75$: 0.00, $s=1$: 0.05; false alarms at $s=0$: 0 of 60. Balanced frame: the $K=30$ agglomerative cut of the WordNet distance matrix (complete linkage) with the smallest variance of cluster sizes, chosen and logged before any depth run; sizes 1–154 classes against 1–316 for the frame of record.

		CIFAR-100 ($K=20$, $n=100$: depth unvalidated)			ImageNet ($K=30$, $n=1000$: depth validated)		
model	input	exc. (r)	$\hat{\delta}_{99.9}$ nat./Eucl.	depth z	exc. (r)	$\hat{\delta}_{99.9}$ nat./Eucl.	depth z
MERU ViT-S	images	-0.020 (199)	0.117/0.117	-1.5	-0.008° (186)	0.088/0.088	-1.5
CLIP ViT-S	images	-0.028 (200)	0.106/0.107	-3.9	-0.006° (171)	0.110/0.088	-1.8
MERU ViT-B	images	-0.020 (200)	0.114/0.114	-2.2	-0.010 (196)	0.087/0.087	-0.9
CLIP ViT-B	images	-0.030 (200)	0.105/0.103	-3.3	-0.009° (187)	0.107/0.085	-1.0
MERU ViT-L	images	-0.012° (195)	0.118/0.118	-0.3	-0.007° (184)	0.088/0.088	-0.5
CLIP ViT-L	images	-0.016 (199)	0.123/0.112	-1.7	-0.010 (197)	0.104/0.084	-1.2
MERU ViT-S	prompts	-0.022 (200)	0.086/0.086	-0.7	-0.003° (155)	0.085/0.086	+0.0
CLIP ViT-S	prompts	-0.023 (200)	0.085/0.082	-0.8	+0.006° (15)	0.093/0.087	+0.8
MERU ViT-B	prompts	-0.012 (200)	0.096/0.096	-0.7	-0.001° (132)	0.089/0.090	+0.8
CLIP ViT-B	prompts	-0.017 (200)	0.094/0.088	+0.3	+0.008° (2)	0.100/0.092	+0.2
MERU ViT-L	prompts	-0.006° (190)	0.102/0.102	+0.6	+0.001° (75)	0.093/0.093	+0.5
CLIP ViT-L	prompts	-0.012 (200)	0.098/0.092	+0.5	+0.006° (5)	0.101/0.092	+0.4

Table 11: **Imposing the geometry does not create depth: MERU and its Euclidean twin under the census of record.** MERU (Desai et al., 2023) embeds images and text on the Lorentz hyperboloid with an entailment objective; its released Euclidean twin is a CLIP baseline with the same ViT, RedCaps data and recipe minus the hyperbolic lift and entailment loss. Both pass through the instrument unchanged: class centroids of the projected embeddings (MERU: space-like hyperboloid coordinates; CLIP: unit vectors), Euclidean distances; exc. (r): excess of $\hat{\delta}_{99.9}$ over the Haar spectrum-matched null with the number of 200 replicates above the real value (°: not genuine under Benjamini–Hochberg over the 24 cells); nat./Eucl.: the same statistic in the model’s own metric (Lorentz distance between tangent-space-mean centroids for MERU, angular for CLIP) against the Euclidean one; depth z : the anisotropic matched-star test of Table 8 (10 star seeds; negative = more hierarchical above the frame than a star), validated at $n=1000$ only (Table 9). On ImageNet images, the validated regime, the excess runs from -0.007 to -0.010 for MERU and from -0.006 to -0.010 for its twin across ViT-S/B/L; MERU’s Lorentz-metric $\hat{\delta}_{99.9}$ differs from its Euclidean value by at most 0.0002; depth z runs from -1.5 to -0.5 for MERU and from -1.8 to -1.0 for CLIP, so neither is more hierarchical than a matched star and MERU is never ahead of its twin. MERU’s hierarchy is a generic $\rightarrow$ specific (text $\supset$ image) partial order, not a class taxonomy.

A trained control, inconclusive. We fine-tuned the census ViT-B/16 for two passes over the images behind its ImageNet centroids with two objectives on the same batches: cross-entropy alone, and cross-entropy plus a hierarchical cross-entropy over the WordNet superclasses. Both fine-tuned models keep the census excess of the frozen backbone (-0.015 and -0.016 against -0.016) and both are certified by the depth test, with z of -2.33 and -2.77 against -3.62 for the frozen backbone: the hierarchical term lowers z relative to plain cross-entropy but neither reaches the frozen reading, so the control neither confirms nor refutes the test (Table 12). The full training set could not be read at training speed from the archive, so the schedule is two passes over the census subset.

ViT-B/16	$\hat{\delta}_{99.9}$	excess	$r/200$	depth	z
frozen (census)	0.054	-0.0162	200	-0.0335	-3.62
cross-entropy	0.048	-0.0153	200	-0.0303	-2.33
cross-entropy + hierarchical CE (WordNet-30)	0.048	-0.0162	200	-0.0295	-2.77

Table 12: **A trained positive control, inconclusive.** The census ViT-B/16 fine-tuned for two passes over the 100 training images per class behind its ImageNet centroids (patch embedding frozen, AdamW, same batches and seed for both objectives), with plain cross-entropy or with an added hierarchical cross-entropy over the WordNet 30-cut; centroids of the same images, census of record and depth test as in Table 8. Both fine-tuned models stay certified and neither reaches the frozen reading; the hierarchical term lowers z relative to plain cross-entropy by less than one unit.

A.6 TRAINING INTERVENTIONS

Clustered structure tracks training. Two interventions move it (Figure 8; raw δ , within architecture): fine-tuning on a task with no class hierarchy raises δ by 19–91%, and δ falls across transformer depth. Shuffling the embedding-to-label correspondence leaves it unchanged, as a null-invariant

quantity should. Random-weight networks give higher raw δ but near-Gaussian features, so they are not counted as evidence; excess-based versions of these interventions await stored point clouds. On matched architecture (ViT-B supervised/DINO/CLIP) all three objectives yield genuine form on most datasets, and the family separations emerge with scale rather than at fixed size.




Figure 8: **The raw reading is learned: training interventions move it.** Raw δ , within architecture. (a) Non-hierarchical fine-tuning raises δ 19–91%. (b) δ falls across transformer depth (layers ≥ 1).

A.7 TEXT CENSUS

Extraction, anisotropy and templates. Extraction is part of the measurement: left-padded batching moves GPT-2’s $\hat{\delta}$ by up to 0.02, of the order of the fragile excesses; the padding-free protocol is the record (Table 15). Genuineness is not a proxy for anisotropy: GTE-base has the highest effective rank and is not genuine. The only LLM-backed embedder tested (GTE-Qwen2-1.5B) is not genuine under the record (Table 14). Causal-LM values are template-sensitive but the scale ordering is template-robust (Tables 39 and 40).




Figure 9: **In text, clustered structure depends on recipe and scale: robust at the larger GPT-2 sizes, fragile at the smaller.** Text census of record (padding-free extraction, Haar null, 99.9th-percentile statistic, 200 replicates; error bars combine null and estimator s.d.; hollow bars are not genuine, BH $p > 0.05$; OLMo-7B, not re-extracted, omitted). Pythia and OLMo-1B are genuine throughout; classic sentence embedders are not on this probe despite raw values equal to genuinely structured vision models.

model	type	d	δ_{norm}	null	excess	z
GPT-2 S	causal LM	768	0.131 ± 0.007	0.083 ± 0.002	+0.048	+6.5
GPT-2 M	causal LM	1024	0.091 ± 0.004	0.103 ± 0.002	-0.013	-2.7
GPT-2 L	causal LM	1280	0.095 ± 0.004	0.129 ± 0.006	-0.034	-5.0
GPT-2 XL	causal LM	1600	0.083 ± 0.005	0.131 ± 0.007	-0.048	-5.4
Pythia-410M	causal LM	1024	0.098 ± 0.005	0.136 ± 0.005	-0.037	-5.4
Pythia-1B	causal LM	2048	0.096 ± 0.004	0.131 ± 0.004	-0.035	-6.2
Pythia-2.8B	causal LM	2560	0.098 ± 0.004	0.139 ± 0.008	-0.041	-4.6
OLMo-1B	causal LM	2048	0.100 ± 0.005	0.101 ± 0.002	-0.001	-0.2
OLMo-7B	causal LM	4096	0.093 ± 0.005	0.108 ± 0.002	-0.015	-2.8
BGE-base	embedder	768	0.121 ± 0.006	0.117 ± 0.005	+0.005	+0.6
BGE-large	embedder	1024	0.121 ± 0.008	0.120 ± 0.003	+0.001	+0.1
GTE-base	embedder	768	0.124 ± 0.011	0.113 ± 0.003	+0.011	+0.9
GTE-large	embedder	1024	0.120 ± 0.012	0.116 ± 0.003	+0.004	+0.4
E5-base	embedder	768	0.115 ± 0.007	0.112 ± 0.005	+0.003	+0.3
E5-large	embedder	1024	0.114 ± 0.010	0.113 ± 0.003	+0.001	+0.1
GTE-Qwen2-1.5B	embedder	1536	0.107 ± 0.004	0.126 ± 0.004	-0.019	-3.1

Table 13: Text models under the calibrated protocol (1000 ImageNet class prompts): δ_{norm} , spectrum-matched null (3 reps) and excess.

model	padding-free extraction, record (Haar $\times$ p99.9, 200 replicates)					original extraction (Table 13)	
	$\delta_{99.9}$	excess	$r/200$	p	2×2	δ	excess
GPT-2 S	0.077	-0.013	197	.020	●○○●	0.131	+0.048
GPT-2 M	0.052	-0.005	170	.154	○○○○	0.091	-0.013
GPT-2 L	0.053	-0.023	200	.005	●●●●	0.095	-0.034
GPT-2 XL	0.048	-0.027	200	.005	●●●●	0.083	-0.048
Pythia-410M	0.067	-0.018	200	.005	●●●●	0.098	-0.037
Pythia-1B	0.066	-0.015	199	.010	●●●●	0.096	-0.035
Pythia-2.8B	0.066	-0.017	200	.005	●●●●	0.098	-0.041
OLMo-1B	0.051	-0.024	200	.005	●●●●	0.100	-0.001
OLMo-7B	—	—	—	—	—	0.093	-0.015
BGE-base	0.078	+0.003	26	.871	○○○○	0.121	+0.005
BGE-large	0.078	+0.002	49	.756	○○○○	0.121	+0.001
GTE-base	0.076	+0.004	11	.945	○○○○	0.124	+0.011
GTE-large	0.075	+0.003	36	.821	○○○○	0.120	+0.004
GTE-Qwen2-1.5B	0.078	+0.002	58	.711	○○○○	0.107	-0.019
E5-base	0.074	+0.004	15	.925	○○○○	0.115	+0.003
E5-large	0.071	+0.002	44	.781	○○○○	0.114	+0.001

Table 14: **The text census of record: padding-free extraction, Haar null, 99.9th-percentile statistic, 200 replicates.** r = replicates above the real value, p the left-tail add-one p -value; genuine = BH-corrected $p \leq 0.05$ over the 15 models. The 2×2 column gives the verdicts under Haar $\times$ p99.9, Haar $\times$ supremum, Gaussian $\times$ p99.9, Gaussian $\times$ supremum (● genuine). OLMo-7B not re-run (exceeds the local GPU). The original extraction (left-padded batches, supremum statistic) is kept for comparison and is not directly comparable. Sign agreement with the 3-replicate original: 13/15. Three verdicts change between extractions, each with a known cause: GPT-2 S (above null in the left-padded original; Table 15), OLMo-1B (the original used a different checkpoint revision), and GTE-Qwen2 (genuine in the original with 3 replicates, -0.019; here excess +0.002, $r = 58$, $p = .711$; not used as evidence).

batch size	GPT-2 M		OLMo-1B	
	fp32 $\hat{\delta}$	fp16 $\hat{\delta}$	fp32 $\hat{\delta}$	fp16 $\hat{\delta}$
1 (no padding)	0.0816	0.0817	0.0817	0.0817
16	0.0996	0.0997	0.0817	0.0816
32	0.0905	0.0905	0.0817	0.0817

Table 15: Where extraction-to-extraction variability comes from. Same 1000 prompts, same model; only batch size and precision vary. Precision is irrelevant (fp16 equals fp32 to four decimals). For GPT-2, left-padded batching changes the hidden states (mean cosine to the padding-free extraction 0.98) and moves $\hat{\delta}$ by up to 0.018; OLMo handles left padding correctly and is invariant. Padding-free (batch-size-1) extraction is therefore the protocol of Table 14.

A.8 CLASS-COUNT VS HIERARCHY DEPTH

Hierarchy depth, not class count. Because C sets the quadruple population, we sweep ImageNet subsets $C \in \{10, \dots, 1000\}$ in two compositions: random (spanning the full hierarchy) and WordNet-coherent (siblings, effectively *flat*). At fixed C , random subsets are more exploitable than coherent ones and carry a larger excess at small C (CLIP-L at every C , DINOv2-L up to $C=100$, DINOv2-G up to $C=20$; beyond that the two compositions converge; Table 16). Random subsets still score -0.08 to -0.14 at $C=10$, so MNIST’s flat verdict is not a class-count artifact, and DINOv2-G stays below its null at every C (excess -0.008 at $C=1000$, $r=200$), where the supremum had placed it at the null.

model	C	random subsets (span the hierarchy)				WordNet-coherent (siblings)			
		$\hat{\delta}_{99.9}$	excess	below	NC adv	$\hat{\delta}_{99.9}$	excess	below	NC adv
DINOv2-L	10	0.030	-0.134	5/5	+1.04	0.068	-0.098	5/5	+0.08
DINOv2-L	20	0.034	-0.091	5/5	+1.32	0.071	-0.065	5/5	+0.12
DINOv2-L	50	0.032	-0.046	5/5	+1.58	0.050	-0.042	5/5	+0.53
DINOv2-L	100	0.031	-0.032	5/5	+1.46	0.044	-0.031	5/5	+0.77
DINOv2-L	200	0.031	-0.022	5/5	+1.45	0.041	-0.023	5/5	+0.59
DINOv2-L	500	0.030	-0.016	2/2	+1.01	0.031	-0.018	2/2	+0.82
DINOv2-L	1000	0.030	-0.014	1/1	+0.57	—	—	—	—
DINOv2-G	10	0.029	-0.136	5/5	+0.80	0.072	-0.095	5/5	-0.56
DINOv2-G	20	0.037	-0.089	5/5	+1.30	0.070	-0.067	5/5	-0.18
DINOv2-G	50	0.039	-0.039	5/5	+1.51	0.053	-0.042	5/5	+0.26
DINOv2-G	100	0.041	-0.022	5/5	+1.30	0.049	-0.028	5/5	+0.36
DINOv2-G	200	0.036	-0.015	5/5	+1.28	0.045	-0.020	5/5	+0.27
DINOv2-G	500	0.034	-0.009	2/2	+0.49	0.034	-0.013	2/2	+0.22
DINOv2-G	1000	0.033	-0.008	1/1	-0.13	—	—	—	—
CLIP-L	10	0.071	-0.080	5/5	-0.08	0.099	-0.066	4/5	-0.36
CLIP-L	20	0.078	-0.047	5/5	+0.20	0.113	-0.036	4/5	-0.24
CLIP-L	50	0.080	-0.027	5/5	+0.34	0.113	-0.014	4/5	-0.46
CLIP-L	100	0.079	-0.024	5/5	+0.21	0.095	-0.016	3/5	-0.31
CLIP-L	200	0.078	-0.019	5/5	+0.00	0.086	-0.010	2/5	-0.35
CLIP-L	500	0.077	-0.012	2/2	-0.27	0.081	-0.011	1/2	-0.42
CLIP-L	1000	0.076	-0.009	1/1	-0.53	—	—	—	—

Table 16: Class-count vs hierarchy-depth deconfound on ImageNet under the census of record (Haar spectrum-matched null, 99.9th-percentile statistic, 200 replicates; means over subset seeds: 5 for $C \leq 200$, 2 for $C=500$, 1 for $C=1000$). “below” = subset seeds whose real value lies below the null at uncorrected $p \leq 0.05$; NC adv = prototype-classifier advantage of the Poincaré readout in pp, from the unchanged task pipeline.

A.9 TREE-MAP METRIC AND LINKAGE CONTROLS

Threshold sensitivity of the selection criterion of §5: the selected configuration is stable for degeneracy bounds in $[0.45, 0.50]$ on ImageNet and $[0.40, 0.50]$ on CIFAR-100; tighter bounds select

cosine-complete on ImageNet (gap 0.13) and cosine-Ward on CIFAR-100 (gap 0.56), and in every case the label-free gap is moderate and configuration-sensitive, never a categorical island. The criterion must be applied per dataset: the configuration it selects on DBpedia (Euclidean, average linkage, where nothing is degenerate) is exactly the degenerate naive configuration on ImageNet, so a criterion transferred across datasets would reintroduce the island.




Figure 10: **The island is an artifact on CIFAR-100 too.** Same reading as Figure 5 on the 20 CIFAR-100 superclass cuts: naive (a) vs criterion-selected (b).

dataset	configuration	max frac.	CPCC	Dv2-B/L/G vs block	Dv2-S+DINO-B vs block	within block
imagenet	euclid-average (deg.)	0.97	0.761	0.03	0.08	0.64
imagenet	euclid-complete (deg.)	0.83	0.606	0.03	0.17	0.44
imagenet	euclid-ward	0.48	0.464	0.15	0.38	0.53
imagenet	cosine-average (selected)	0.40	0.691	0.38	0.35	0.48
imagenet	cosine-complete	0.32	0.539	0.13	0.17	0.50
imagenet	cosine-ward	0.37	0.464	0.17	0.39	0.55
cifar100	euclid-average (deg.)	0.80	0.798	0.00	0.41	0.57
cifar100	euclid-complete (deg.)	0.66	0.722	0.17	0.40	0.54
cifar100	euclid-ward	0.33	0.534	0.37	0.65	0.63
cifar100	cosine-average (deg.)	0.53	0.777	0.30	0.54	0.58
cifar100	cosine-complete (selected)	0.39	0.698	0.39	0.51	0.63
cifar100	cosine-ward	0.14	0.584	0.56	0.62	0.61

Table 17: Tree-map configurations with the two quantities that drive the selection criterion of §5: the degeneracy diagnostic (largest-cluster fraction at the reference cut, worst model) and the mean cophenetic fidelity (CPCC) of each model’s dendrogram to its own distance matrix. Configurations with max-cluster fraction > 0.5 are excluded; among the rest, highest CPCC selects the marked row. Mean pairwise ARIs at the reference cut shown for all configurations. Under the naive ImageNet configuration, cross-family ARI between the supervised ViTs and the contrastive VLMs reaches 0.82 while every DINOv2-vs-block pair stays at or below 0.06, and the worst model’s largest cluster holds 97% of the classes: the island is made by chaining.

dataset	configuration	ARI at the cut	cophenetic corr.	triplet agreement
ImageNet	euclid-average	0.03 / 0.64	0.36 / 0.80	0.47 / 0.74
ImageNet	euclid-complete	0.03 / 0.44	0.20 / 0.66	0.41 / 0.65
ImageNet	euclid-ward	0.16 / 0.53	0.62 / 0.84	0.67 / 0.85
ImageNet	cosine-average	0.39 / 0.48	0.48 / 0.78	0.77 / 0.76
ImageNet	cosine-complete	0.13 / 0.49	0.18 / 0.72	0.36 / 0.71
ImageNet	cosine-ward	0.17 / 0.55	0.59 / 0.85	0.69 / 0.86
CIFAR-100	euclid-average	0.00 / 0.57	0.15 / 0.78	0.41 / 0.74
CIFAR-100	euclid-complete	0.17 / 0.54	0.23 / 0.72	0.41 / 0.66
CIFAR-100	euclid-ward	0.37 / 0.63	0.63 / 0.77	0.70 / 0.79
CIFAR-100	cosine-average	0.30 / 0.58	0.42 / 0.72	0.51 / 0.64
CIFAR-100	cosine-complete	0.39 / 0.63	0.46 / 0.64	0.50 / 0.61
CIFAR-100	cosine-ward	0.56 / 0.61	0.74 / 0.73	0.76 / 0.76

Table 18: **The island is a property of the cut and of the Euclidean configurations.** Each entry: DINOv2-B/L/G vs the supervised+contrastive block / within the block. ARI between cuts at 30 (ImageNet) or 20 (CIFAR-100) clusters; Pearson correlation between the two trees’ cophenetic distance vectors; agreement on which pair merges first over 10^4 random class triplets. Dendrograms rebuilt from the census cache with the six configurations of Table 17.

model	euclid-avg (naive, degen.)	euclid-ward	cosine-avg (degen. C100)	cosine-complete ← sel.	cosine-ward
ViT-T	0.31	0.38	0.29	0.33	0.33
ViT-S	0.35	0.56	0.37	0.47	0.56
ViT-B	0.39	0.61	0.43	0.45	0.57
ViT-L	0.38	0.55	0.49	0.48	0.60
DINO-B	0.22	0.48	0.28	0.38	0.54
DINOv2-S	0.12	0.51	0.32	0.33	0.49
DINOv2-B	0.02	0.47	0.35	0.42	0.51
DINOv2-L	0.01	0.32	0.25	0.50	0.48
DINOv2-G	0.00	0.19	0.08	0.20	0.52
CLIP-B	0.18	0.48	0.18	0.37	0.48
CLIP-L	0.35	0.52	0.32	0.41	0.54
SigLIP-B	0.22	0.49	0.22	0.37	0.52

Table 19: Recovery of the true CIFAR-100 superclasses (ARI of the 20-cluster cut) per model and configuration, the quantity behind the corrected picture of §5. Degenerate configurations marked; the criterion-selected configuration is cosine-complete. Paired within each of the 3 admissible configurations, the label-supervised ViT-B recovers more than each DINOv2 model in 11 of 12 comparisons; DINOv2 itself recovers the superclasses at ARI 0.19–0.52.

A.10 THE COROLLARY: DIAGNOSING EXPLOITABILITY

One parameter-free projection and two tasks serve the corollary. The parameter-free map $\phi(x) = \tanh(\frac{1}{2}\|s(x - \mu)\|) \cdot \frac{s(x - \mu)}{\|s(x - \mu)\|}$ centers features, scales the 95th-percentile norm to radius $1/\sqrt{2}$, and preserves directions; prototypes in the ball are means of projected points re-mapped through the origin. Tasks: nearest-centroid (NC) and 5-way 5-shot (FS, 1000 paired episodes); metric arms: Euclidean (R), Poincaré (H), the radial map with Euclidean distances (RT), cosine (COS), and the L2-normalized stack. Changing the readout metric on the intact configuration, which is what the rule selects among, is distinct from deforming the configuration itself.

δ_{norm} predicts the gain within datasets. Within each dataset the Poincaré advantage tracks δ_{norm} (NC $r = -0.83$ on ImageNet) and survives partialing out dimension, baseline accuracy, family and spectral covariates; without DINOv2 it holds on ImageNet and CIFAR-100 but not on CIFAR-10/DTD (Appendix A.13). Pooled cross-dataset correlations are family-driven and not used as evidence; on flat-label datasets no metric helps. Gains carry paired confidence intervals and a McNemar test (Appendix A.13).

The training objective picks the metric. The radial map with Euclidean distances changes nothing, so the gain is the metric itself, and the split follows the objective: cosine suffices for DINOv2,

Poincaré $\approx$ cosine for supervised ViTs, and Poincaré beats cosine for CLIP by up to +1.3pp even after L2 normalization (Figure 11 in Appendix A.13; held-out selection and policy comparison in this appendix). The rule is read off the families it is scored on: it describes them rather than validating a policy.

A diagnostic, not a policy. Khrulkov et al. (2020) estimate $c \approx 0.33$ from the raw reading of their datasets; the same rule on structureless Gaussian clouds gives the curvatures of Section 6. Cosine is a safe default, gating on δ_{norm} is the weakest policy tested (this appendix), and gains are modest and limited to prototype-based tasks (sample-level tasks never benefit, Appendix A.18). On DBpedia, a three-level text hierarchy, the diagnostic predicts correctly on both sides: genuine excess, no metric gain.

A.11 TASK ACCURACIES (REGENERATED)

model	IN	C100	C10	DTD	FMNIST	MNIST
ViT-T	68.1/67.2	29.3/28.2	87.5/87.4	63.7/62.5	78.5/78.4	75.9/74.0
ViT-S	77.7/77.5	74.9/74.6	93.4/93.2	69.3/68.5	80.7/80.4	84.8/84.1
ViT-B	81.6/81.7	77.3/77.0	94.5/94.0	70.8/70.8	81.3/81.6	81.6/80.5
ViT-L	86.1/86.2	85.3/85.2	97.5/97.5	71.8/71.5	83.7/83.5	84.2/83.6
DINOv2-S	74.7/74.4	73.4/74.4	95.0/95.2	73.4/73.4	82.2/82.2	83.4/81.9
DINOv2-B	78.9/78.9	87.4/87.9	97.1/97.4	74.2/74.7	83.2/82.8	80.6/78.8
DINOv2-L	79.8/80.4	85.6/87.7	97.9/98.3	74.4/75.2	84.4/84.5	74.1/72.7
DINOv2-G	78.4/78.3	89.8/91.0	97.8/98.2	72.9/73.1	85.9/85.8	79.3/78.6
CLIP-B	61.4/59.9	65.6/65.1	87.9/88.3	63.7/63.2	75.9/75.1	70.1/66.6
CLIP-L	75.0/74.5	74.6/74.5	95.9/95.9	68.6/69.0	76.2/75.9	85.2/86.1

Table 20: Nearest-centroid accuracy (%), Euclidean/Poincaré, regenerated from the audited reruns.

model	IN	C100	C10	DTD	FMNIST	MNIST
ViT-T	97.5/97.6	62.4/63.1	87.6/88.1	83.3/83.9	81.6/81.9	74.3/74.6
ViT-S	98.6/98.7	92.2/92.4	92.0/92.5	85.4/86.0	83.7/83.5	80.5/81.4
ViT-B	98.9/99.0	92.5/92.7	92.7/93.0	85.4/86.2	83.4/83.3	76.8/77.6
ViT-L	99.4/99.4	94.8/95.5	95.0/95.8	85.6/86.5	84.4/84.7	80.8/81.6
DINOv2-S	97.4/97.9	94.3/95.1	92.9/93.5	89.8/90.5	85.4/85.5	80.2/80.8
DINOv2-B	97.0/97.8	96.5/97.2	94.7/95.4	89.2/90.3	85.9/85.9	76.6/77.1
DINOv2-L	96.2/97.1	96.1/97.1	94.1/95.1	87.1/88.4	87.8/87.9	70.5/71.1
DINOv2-G	96.2/97.1	95.9/97.2	92.3/93.8	87.2/88.5	89.0/89.2	72.9/73.8
CLIP-B	96.8/96.8	87.7/88.9	89.5/90.1	82.2/83.1	78.8/78.9	70.9/69.7
CLIP-L	98.4/98.6	90.0/90.8	92.4/93.0	85.2/86.2	76.8/77.3	83.7/84.7

Table 21: 5-way 5-shot accuracy (%), Euclidean/Poincaré, 1000 paired episodes.

A.12 METRIC CONTROLS

model	IN	C100	C10	DTD	FMNIST	MNIST
ViT-T	-0.01	+1.01	+0.66	-0.23	+0.21	+0.64
ViT-S	+0.17	+0.59	+0.56	-0.03	+0.28	+0.34
ViT-B	+0.20	+0.56	+0.49	+0.28	+0.03	+0.19
ViT-L	+0.11	+1.17	+1.25	+0.13	+0.00	+0.08
DINOv2-S	+0.68	+0.80	-0.36	+0.30	+0.02	-0.07
DINOv2-B	+1.40	+1.11	+0.27	+1.48	+0.21	+0.17
DINOv2-L	+2.13	+1.80	+0.92	+2.40	+0.29	+0.25
DINOv2-G	+2.57	+2.44	+2.51	+2.63	+0.38	+0.30
CLIP-B	-0.04	-0.03	-0.19	-0.23	+0.06	+0.04
CLIP-L	+0.01	+0.12	+0.16	-0.08	+0.12	+0.08

Table 22: Few-shot advantage of cosine over Euclidean (pp).

model	IN	C100	C10	DTD	FMNIST	MNIST
ViT-T	+0.01	+0.45	+0.19	+0.32	+0.10	+0.52
ViT-S	-0.02	+0.16	+0.27	+0.25	+0.08	+0.65
ViT-B	-0.04	+0.12	+0.18	+0.19	+0.09	+0.51
ViT-L	-0.03	+0.31	+0.35	+0.16	+0.17	+0.42
DINOv2-S	+0.03	+0.15	+0.13	+0.04	+0.07	+0.45
DINOv2-B	+0.04	+0.09	+0.03	+0.01	+0.04	+0.48
DINOv2-L	+0.02	+0.09	-0.00	+0.03	+0.03	+0.43
DINOv2-G	+0.03	+0.13	-0.02	+0.01	+0.04	+0.56
CLIP-B	-0.00	+0.44	+0.20	+0.36	+0.17	+0.23
CLIP-L	+0.00	+0.39	+0.25	+0.36	+0.33	+0.58

Table 23: Few-shot advantage of the radial-rescaling control over Euclidean (pp): the map without the metric changes nothing.

A.13 STATISTICS




Figure 11: **More tree-like cells carry larger zero-cost metric gains.** Best zero-cost metric minus Euclidean per cell vs $\hat{\delta}_{\text{norm}}$ (FS left, NC right). Within-dataset correlations -0.60 to -0.83 (FS) and -0.61 to -0.86 (NC); pooled correlations are family-driven and not used as evidence (Appendix A.13).

measure	raw	null mean	$\tau_{0.05}$	calibrated	null-centered	pairs $p < 0.05$
mutual-kNN ($k=10$), Euclidean	0.472	0.010	0.015	0.464	0.462	66/66
mutual-kNN, Poincaré	0.425	0.010	0.019	0.414	0.415	66/66
linear CKA, Euclidean	0.633	0.108	0.111	0.577	0.524	66/66
linear CKA, Poincaré	0.640	0.108	0.111	0.586	0.532	66/66

Table 24: **Cross-model agreement survives permutation calibration.** Means over the 66 model pairs of the 12 vision backbones on the 1000 ImageNet class centroids (census cache). Calibration of Gröger et al. (2026) with $K=200$ permutations of the class correspondence and $\alpha=0.05$, scalar (no layer search): $\tau_{0.05}$ is the $\lceil 0.95(K+1) \rceil$ -th order statistic of the observed score and its nulls (eq. 9), $p = (1 + \#\{\text{null} \geq \text{obs}\}) / (K+1)$ (eq. 10), calibrated = $\max\{(\text{obs} - \tau_{0.05}) / (1 - \tau_{0.05}), 0\}$ (eq. 12); “null-centered” is the earlier raw-minus-null-mean variant, kept for continuity. The mKNN null mean equals the analytic chance level $k/(n-1) = 10/999$.

quantity	dataset	pooled r	demeaned r	ViT	DINOv2	CLIP
NC H-R	ImageNet	-0.83	-0.73	-0.99	-0.77	(-1)
	CIFAR-100	-0.87	-0.45	-0.68	-0.33	(-1)
	CIFAR-10	-0.63	-0.31	+0.01	-0.98	(+1)
	DTD	-0.46	-0.57	-0.78	-0.76	(+1)
FS H-R	ImageNet	-0.57	-0.73	-0.04	-0.97	(-1)
	CIFAR-100	-0.67	-0.63	-0.41	-0.96	(+1)
	CIFAR-10	-0.78	-0.64	-0.22	-0.79	(-1)
	DTD	-0.71	-0.86	-0.95	-0.98	(+1)
FS best-R	ImageNet	-0.64	-0.71	-0.57	-0.90	(-1)
	CIFAR-100	-0.83	-0.82	-0.58	-0.96	(+1)
	CIFAR-10	-0.63	-0.55	-0.26	-0.65	(-1)
	DTD	-0.60	-0.70	-0.95	-0.88	(+1)

Table 25: Family-identity control for the δ_{norm} -gain correlations (within dataset, 10 backbones = 3 families with nested scales). Demeaned: both variables centered within family before correlating; ViT/DINOv2: within-family correlations (4 models each); CLIP has only two models, so its within-family value (parenthesized sign) is not informative. Without DINOv2 (six backbones) the NC correlation holds on ImageNet and CIFAR-100 ($-0.87/-0.65$) but not on CIFAR-10/DTD ($+0.23/-0.34$). The within-dataset prediction survives the family control; pooled cross-dataset correlations of the best-metric advantage do not ($-0.31 \rightarrow -0.18$ over all 60 cells, $-0.33 \rightarrow -0.04$ on hierarchical cells) and are therefore family-driven.

model	IN	C100	C10	DTD	FMNIST	MNIST
ViT-T	$10^{-22}(-)$	$10^{-5}(-)$	0.19(-)	0.03(-)	0.38(-)	$10^{-13}(-)$
ViT-S	0.09(-)	0.10(-)	0.02(-)	0.10(-)	0.08(-)	$10^{-4}(-)$
ViT-B	0.21(+)	0.07(-)	$10^{-5}(-)$	1.00(-)	0.08(+)	$10^{-7}(-)$
ViT-L	0.58(+)	0.59(-)	0.76(+)	0.69(-)	0.12(-)	0.00(-)
DINOv2-S	0.00(-)	$10^{-6}(+)$	0.14(+)	1.00(+)	1.00(-)	$10^{-11}(-)$
DINOv2-B	0.57(+)	$10^{-4}(+)$	$10^{-4}(+)$	0.31(+)	0.01(-)	$10^{-16}(-)$
DINOv2-L	$10^{-14}(+)$	$10^{-38}(+)$	$10^{-8}(+)$	0.11(+)	0.40(+)	$10^{-12}(-)$
DINOv2-G	0.10(-)	$10^{-16}(+)$	$10^{-13}(+)$	0.81(+)	0.72(-)	0.00(-)
CLIP-B	$10^{-40}(-)$	0.07(-)	$10^{-4}(+)$	0.46(-)	$10^{-4}(-)$	$10^{-45}(-)$
CLIP-L	$10^{-8}(-)$	0.66(-)	0.59(-)	0.45(+)	0.09(-)	$10^{-7}(+)$

Table 26: McNemar p -values for NC test-set decisions, Poincaré vs Euclidean (order of magnitude when $p < 10^{-3}$; sign of the NC advantage in parentheses, so significance is readable with its direction).

task	dataset	r	CI95
NC	imagenet	-0.83	[-0.97, -0.65]
NC	cifar100	-0.87	[-0.97, -0.78]
NC	cifar10	-0.63	[-0.97, -0.06]
NC	dtd	-0.46	[-0.90, +0.22]
FS	imagenet	-0.57	[-0.91, +0.14]
FS	cifar100	-0.67	[-0.95, -0.15]
FS	cifar10	-0.78	[-0.93, -0.23]
FS	dtd	-0.71	[-0.94, -0.32]
NC	POOLED(cluster-bs)	-0.45	[-0.60, -0.36]
FS	POOLED(cluster-bs)	-0.18	[-0.50, +0.12]

Table 27: Bootstrap CI95 for the correlations between δ_{norm} and the projection advantage (within-dataset: 10k resamples of the 10 models; pooled: cluster bootstrap over datasets).

policy (test advantage over Euclidean, pp)	all 60 cells	hierarchical (40)
always cosine	+0.23	+0.28
objective rule (zero validation)	+0.28	+0.41
δ -gated objective rule	+0.08	+0.13
validation-picked per cell	+0.63	+0.78
test oracle (reference)	+0.64	+0.78

Table 28: Metric-selection policies evaluated on held-out episodes (500 validation / 500 test per cell, paper protocol). Validation picking retains the oracle advantage; the zero-validation objective rule modestly beats always-cosine.

NC policy (advantage over Euclidean, pp)	all 60	flat (20)	hierarchical (40)
always cosine	+0.48	+0.24	+0.59
always Poincaré	-0.24	-0.70	-0.00
gated rule (flat→R)	+0.34	+0.00	+0.51
δ -gated rule	+0.19	+0.00	+0.29

Table 29: NC metric-selection policies. The gate’s value is avoiding the Poincaré projection on flat-label data; cosine is a safe NC default everywhere.

A.14 PROJECTION-TARGET SENSITIVITY

target t	0.25	0.35	0.41	0.45	0.50	0.58	0.71	0.95
FS adv. mean (pp)	+0.88	+0.87	+0.87	+0.86	+0.84	+0.82	+0.76	+0.56
NC adv. mean (pp)	-0.30	-0.25	-0.23	-0.21	-0.17	-0.12	-0.01	+0.15
% cells FS > 0	97	100	100	100	100	100	100	100

Table 30: Sensitivity to the projection target t (36 hierarchical cells). The headline $t = 1/\sqrt{2} \approx 0.71$ was fixed a priori and is conservative for FS.

A.15 ESTIMATOR ROBUSTNESS

model	data	excess vs quadruple budget ($10^4 \dots 2 \times 10^5$)						p99.9 @ 5×10^5
ViT-L	IN	-0.026	-0.024	-0.027	-0.030	-0.034	-0.035	-0.018
DINOv2-L	IN	-0.027	-0.019	-0.024	-0.011	-0.004	-0.005	-0.024
CLIP-B	IN	-0.008	-0.012	-0.023	-0.025	-0.025	-0.031	-0.007
ViT-L	C100	-0.045	-0.047	-0.047	-0.049	-0.052	-0.053	-0.033
DINOv2-L	C100	-0.071	-0.055	-0.054	-0.050	-0.052	-0.055	-0.069
CLIP-B	C100	-0.045	-0.047	-0.050	-0.049	-0.052	-0.051	-0.030
ViT-L	DTD	-0.064	-0.070	-0.076	-0.083	-0.084	-0.084	-0.054
DINOv2-L	DTD	-0.068	-0.072	-0.075	-0.083	-0.084	-0.084	-0.064
CLIP-B	DTD	-0.048	-0.055	-0.059	-0.065	-0.066	-0.066	-0.035
Centroid bootstrap (30 resamples of the per-class images): s.d. of the excess								
ViT-L	C100	excess	-0.0493	bootstrap s.d.	0.0029	(30 resamples, 500 images/class)		
DINOv2-L	C100	excess	-0.0493	bootstrap s.d.	0.0022	(30 resamples, 500 images/class)		
DINOv2-G	C100	excess	-0.0778	bootstrap s.d.	0.0020	(30 resamples, 500 images/class)		
CLIP-B	C100	excess	-0.0476	bootstrap s.d.	0.0017	(30 resamples, 500 images/class)		
ViT-L	DTD	excess	-0.0803	bootstrap s.d.	0.0049	(30 resamples, 80 images/class)		
DINOv2-L	DTD	excess	-0.0878	bootstrap s.d.	0.0048	(30 resamples, 80 images/class)		
DINOv2-G	DTD	excess	-0.0863	bootstrap s.d.	0.0051	(30 resamples, 80 images/class)		
CLIP-B	DTD	excess	-0.0581	bootstrap s.d.	0.0054	(30 resamples, 80 images/class)		

Table 31: Estimator robustness. Top: on CIFAR-100 and DTD the excess is stable once the quadruple budget reaches 10^5 (the paper uses 5×10^5); on ImageNet ($C=1000$) every sign holds at every budget but magnitudes drift by up to 0.02 (DINOv2-L toward its null, ViT-L and CLIP-B away from it), so ImageNet excess magnitudes carry a budget caveat. The 99.9th-percentile variant of the four-point statistic, robust to the supremum, is negative wherever the supremum excess is. Bottom: resampling the per-class images that form each centroid (reported on CIFAR-100 and DTD) moves the excess by an order of magnitude less than its size.

model	IN			C100			C10			DTD			FMNIST			MNIST		
	exc.	r	p	exc.	r	p	exc.	r	p	exc.	r	p	exc.	r	p	exc.	r	p
ViT-T	-0.005°	169	.159	-0.006°	151	.249	-0.021°	172	.144	-0.014	197	.020	+0.018°	34	.831	-0.026°	168	.164
ViT-S	-0.012	200	.005	-0.000°	111	.448	-0.056	200	.005	-0.019	199	.010	-0.033	196	.025	-0.037°	181	.100
ViT-B	-0.016	200	.005	+0.006°	43	.786	-0.050	200	.005	-0.021	200	.005	-0.042	200	.005	-0.013°	142	.294
ViT-L	-0.016	200	.005	-0.020	200	.005	-0.060	200	.005	-0.031	200	.005	-0.051	200	.005	-0.040°	193	.040
DINO-B	-0.020	200	.005	-0.032	200	.005	-0.061	200	.005	-0.027	200	.005	-0.007°	127	.368	-0.034°	188	.065
DINOv2-S	-0.016	200	.005	-0.031	200	.005	-0.071	200	.005	-0.018	200	.005	-0.031°	193	.040	-0.004°	112	.443
DINOv2-B	-0.017	200	.005	-0.031	200	.005	-0.112	200	.005	-0.035	200	.005	-0.062	200	.005	-0.017°	147	.269
DINOv2-L	-0.014	200	.005	-0.035	200	.005	-0.128	200	.005	-0.042	200	.005	-0.035	197	.020	-0.014°	132	.343
DINOv2-G	-0.008	200	.005	-0.040	200	.005	-0.130	200	.005	-0.039	200	.005	-0.046	199	.010	-0.022°	158	.214
CLIP-B	-0.007°	172	.144	-0.019	200	.005	-0.036	195	.030	-0.024	200	.005	-0.035°	181	.100	-0.001°	104	.483
CLIP-L	-0.009	197	.020	-0.018	200	.005	-0.043	199	.010	-0.027	200	.005	-0.037°	188	.065	-0.008°	126	.373
SigLIP-B	-0.008°	189	.060	-0.022	200	.005	-0.050	199	.010	-0.023	200	.005	-0.042	198	.015	-0.044	198	.015

Table 32: **The census of record per cell: Haar-rotated spectrum-matched null, 99.9th-percentile statistic, 200 replicates.** Excess over the null mean; r = replicates above the real value; p = $(1 + \#\{\text{null} \leq \text{real}\})/201$; genuine = Benjamini–Hochberg-corrected $p \leq 0.05$ over the 72 cells (°: not genuine). Genuine cells: 49/72 overall, 18/24 on ImageNet+CIFAR-100. Sign agreement with the original 5-replicate Gaussian $\times$ supremum census on the 72 comparable cells: 67/72. Every cell on the census cache (one centroid source). The four null $\times$ statistic verdicts per cell are in Table 33.

model	IN	C100	C10	DTD	FMNIST	MNIST
ViT-T	○●○○	○●○○	○○○○	●●●●	○○○○	○○○○
ViT-S	●●●●	○●○○	●●●●	●●●●	●●●●	○○○○
ViT-B	●●●●	○○○○	●●●●	●●●●	●●●●	○○○○
ViT-L	●●●●	●●●●	●●●●	●●●●	●●●●	○○●●
DINO-B	●○○○	●●●●	●●●●	●●●●	○○○○	○○○○
DINOv2-S	●○○○	●●●●	●●●●	●●●●	○○●●	○○○○
DINOv2-B	●○○○	○●○○	●●●●	●●●●	●●●●	○○○○
DINOv2-L	●○○●	○●○○	●●●●	●●●●	●●●●	○○○○
DINOv2-G	●○○○	●●●●	●●●●	●●●●	●●●●	○○○○
CLIP-B	○●○○	●●●●	●●●●	●●●●	○○○○	○○○○
CLIP-L	●●●●	●●●●	●●●●	●●●●	○○○○	○○○○
SigLIP-B	○○●●	●●●●	●●●●	●●●●	●●●●	●●●●

Table 33: **Verdicts under all four null $\times$ statistic constructions, per cell.** Each cell shows four symbols in the order Haar $\times$ p99.9 (the record), Haar $\times$ supremum, Gaussian $\times$ p99.9, Gaussian $\times$ supremum; ● = genuine (BH-corrected $p \leq 0.05$ over the 72 cells of that construction), ○ = not. Genuine counts: 49/72, 46/72, 52/72, 52/72. The Haar construction reproduces the sample spectrum exactly; the Gaussian one does not at small n and inflates excesses for low-rank clouds (Table 34). The supremum is decided by a few extreme quadruples and disagrees with the 99.9th percentile in both directions on DINOv2-S/B/G ImageNet (genuine under both p99.9 constructions, above every replicate under both supremum constructions); the 99.9th percentile is decided by the bulk.

model	IN			C100			C10		
	Gauss.	Haar	PC-perm.	Gauss.	Haar	PC-perm.	Gauss.	Haar	PC-perm.
ViT-T	-0.005	-0.014	-0.038	-0.044	-0.031	-0.055	-0.057	-0.027	-0.050
ViT-S	-0.023	-0.025	-0.052	-0.034	-0.025	-0.049	-0.098	-0.063	-0.090
ViT-B	-0.027	-0.025	-0.054	-0.018	-0.011	-0.039	-0.090	-0.054	-0.078
ViT-L	-0.028	-0.029	-0.040	-0.047	-0.034	-0.051	-0.104	-0.066	-0.096
DINO-B	-0.005	-0.002	-0.008	-0.068	-0.056	-0.068	-0.100	-0.070	-0.097
DINOv2-S	+0.010	+0.015	+0.013	-0.050	-0.029	-0.051	-0.093	-0.083	-0.098
DINOv2-B	+0.004	+0.015	+0.006	-0.042	-0.007	-0.041	-0.132	-0.120	-0.124
DINOv2-L	-0.013	+0.000	-0.011	-0.053	-0.008	-0.044	-0.153	-0.136	-0.122
DINOv2-G	-0.000	+0.016	+0.001	-0.080	-0.032	-0.074	-0.155	-0.138	-0.110
CLIP-B	-0.028	-0.028	-0.048	-0.046	-0.036	-0.053	-0.079	-0.045	-0.075
CLIP-L	-0.016	-0.017	-0.029	-0.042	-0.031	-0.052	-0.087	-0.049	-0.082
SigLIP-B	-0.019	-0.020	-0.044	-0.042	-0.032	-0.046	-0.096	-0.056	-0.081

Table 34: Excess under three constructions of the spectrum-matched null (10 replicates each): Gaussian coefficients (the paper’s), Haar-rotated coefficients (exact sample spectrum) and PC-permutation (exact per-component marginals). The Gaussian and PC-permutation constructions agree; the exact-sample-spectrum construction yields smaller excesses for low-effective-rank clouds on the small- C sets (most visibly DINOv2 on CIFAR-100), so the size of those excesses is partly a property of the null hypothesis, while signs and the ImageNet readings are stable.

model	IN		C100		C10		DTD		FMNIST		MNIST	
	exc.	r	exc.	r	exc.	r	exc.	r	exc.	r	exc.	r
ViT-T	-0.012	200	-0.014 [°]	191	-0.043	198	-0.017	199	+0.007 [°]	73	-0.026 [°]	167
ViT-S	-0.017	200	-0.017	200	-0.073	200	-0.021	200	-0.049	200	-0.036 [°]	183
ViT-B	-0.022	200	-0.006 [°]	170	-0.066	200	-0.021	200	-0.041	199	-0.010 [°]	130
ViT-L	-0.021	200	-0.024	200	-0.080	200	-0.024	200	-0.066	200	-0.043	196
DINO-B	-0.021	200	-0.035	200	-0.070	200	-0.024	200	-0.006 [°]	116	-0.038 [°]	193
DINOv2-S	-0.014	200	-0.033	200	-0.084	200	-0.018	200	-0.054	200	-0.004 [°]	113
DINOv2-B	-0.013	200	-0.032	200	-0.140	200	-0.021	200	-0.081	200	-0.019 [°]	147
DINOv2-L	-0.010	200	-0.029	200	-0.151	200	-0.024	200	-0.051	200	-0.020 [°]	145
DINOv2-G	-0.006	200	-0.029	200	-0.169	200	-0.017	200	-0.064	200	-0.023 [°]	157
CLIP-B	-0.017	200	-0.023	200	-0.039	197	-0.031	200	-0.038 [°]	185	+0.002 [°]	93
CLIP-L	-0.017	200	-0.025	200	-0.045	199	-0.031	200	-0.038 [°]	192	-0.008 [°]	124
SigLIP-B	-0.013	200	-0.025	200	-0.053	200	-0.027	200	-0.052	199	-0.044	197

Table 35: **The angular reading agrees with the Euclidean census.** Centroids L2-normalized, spherical geodesic distances, Haar null built on the normalized cloud and re-normalized, 99.9th-percentile statistic, 200 replicates, BH over 72 cells ([°]: not genuine); a robustness reading, not the record. Sign-negative 70/72; genuine 56/72, 22/24 on ImageNet+CIFAR-100; verdict agreement with the Euclidean record 65/72 cells. Text (same protocol on the padding-free embeddings): genuine 8/15: gpt2-m, gpt2-l, gpt2-xl, pythia-410m, pythia-1b, pythia-2b8, olmo-1b, gte-qwen2. The angular tree: sibling-triplet agreement on CIFAR-100 for DINOv2-L is 0.91 under cosine against 0.71 under Euclidean distance.

model	excess (reference)	bootstrap mean	bootstrap s.d.	fraction negative
ViT-T	-0.0055	-0.0051	0.0011	1.00
ViT-S	-0.0126	-0.0122	0.0007	1.00
ViT-B	-0.0165	-0.0160	0.0006	1.00
ViT-L	-0.0160	-0.0160	0.0004	1.00
DINO-B	-0.0209	-0.0208	0.0003	1.00
DINOv2-S	-0.0164	-0.0162	0.0004	1.00
DINOv2-B	-0.0167	-0.0164	0.0002	1.00
DINOv2-L	-0.0136	-0.0134	0.0001	1.00
DINOv2-G	-0.0080	-0.0083	0.0004	1.00
CLIP-B	-0.0068	-0.0069	0.0009	1.00
CLIP-L	-0.0094	-0.0095	0.0006	1.00
SigLIP-B	-0.0080	-0.0081	0.0006	1.00

Table 36: **The ImageNet excess is stable under resampling of the 100 training images per class.** 30 bootstrap resamples per backbone under the record protocol (Haar null, 99.9th-percentile statistic; 20 null replicates per resample). The bootstrap s.d. of the excess is at most 0.0011; every resample of every backbone is sign-negative.

model	IN	C100	C10	DTD	FMNIST	MNIST
ViT-T	-1.1 (0) ^o	-0.7 (0) ^o	-1.1 (0) ^o	-2.0 (13) [†]	+0.9 (0) ^o	-1.0 (0) ^o
ViT-S	-3.6 (30)	-0.0 (0) ^o	-2.7 (30)	-3.1 (30)	-1.7 (30) [†]	-1.3 (0) ^o
ViT-B	-5.2 (30)	+0.8 (0) ^o	-2.4 (30)	-4.0 (30)	-2.1 (30)	-0.6 (0) ^o
ViT-L	-5.0 (30)	-4.0 (30)	-2.8 (30)	-6.4 (30)	-2.5 (30)	-1.6 (0) ^o
DINO-B	-8.9 (30)	-7.6 (30)	-2.7 (30)	-4.5 (30)	-0.3 (0) ^o	-1.5 (0) ^o
DINOv2-S	-15.9 (30)	-9.7 (30)	-2.8 (30)	-3.2 (30)	-1.6 (30) ^o ‡	-0.2 (0) ^o
DINOv2-B	-31.6 (30)	-14.7 (30)	-4.3 (30)	-8.1 (30)	-2.7 (30)	-0.6 (0) ^o
DINOv2-L	-30.7 (30)	-21.6 (30)	-4.9 (30)	-10.6 (30)	-1.8 (29) [†]	-0.5 (0) ^o
DINOv2-G	-11.5 (30)	-28.5 (30)	-4.8 (30)	-8.2 (30)	-2.0 (30)	-0.8 (0) ^o
CLIP-B	-1.3 (0) ^o	-3.1 (30)	-1.8 (7) [†]	-3.5 (30)	-1.3 (0) ^o	-0.0 (0) ^o
CLIP-L	-2.1 (12) [†]	-3.5 (30)	-2.0 (30) [†]	-4.3 (30)	-1.5 (0) ^o	-0.4 (0) ^o
SigLIP-B	-1.7 (0) ^o	-4.3 (30)	-2.6 (30)	-4.1 (30)	-2.0 (30) [†]	-1.9 (30) [†]

Table 37: **The genuine count survives centroid resampling and estimator noise.** Per cell: $z_{\text{joint}} = \text{excess} / \sqrt{\sigma_{\text{null}}^2 + \sigma_{\text{boot}}^2 + \sigma_{\text{est}}^2}$ (200-replicate null s.d., bootstrap s.d. of the excess over 30 centroid resamples with replacement of the cached images per class, quadruple-seed s.d.), and in parentheses the number of the 30 resamples in which the cell is genuine under Benjamini–Hochberg over the 72 cells (50 Haar replicates per resample). ^o: not genuine in the census of record; [†]: record-genuine but not under one of the two criteria ($z_{\text{joint}} \leq -2$, genuine in ≥ 27 of 30); [‡]: not record-genuine but passing one of them. Counts: record 49/72 (18/24); $z_{\text{joint}} \leq -2$: 42/72 (18/24); genuine in ≥ 27 of 30 resamples: 47/72 (17/24).

A.16 ORC EDGE TYPES

High mean ORC coexists with tree-likeness because they live on different edges: negative curvature concentrates on between-superclass bridges (about twice as often as within superclasses on ImageNet; Table 38), while the mean measures within-cluster compactness. DINOv2 on ImageNet inverts the pattern, consistent with its at-null supremum reading, and recovers it on CIFAR-100.

model	dataset	ORC within	ORC across	%neg within	%neg across
ViT-T	IN	0.35	0.27	5.2	9.4
ViT-T	C100	0.51	0.42	1.4	2.7
ViT-S	IN	0.37	0.28	3.4	8.3
ViT-S	C100	0.61	0.48	0.0	3.0
ViT-B	IN	0.40	0.30	2.1	5.9
ViT-B	C100	0.61	0.48	0.0	1.7
ViT-L	IN	0.40	0.31	2.1	6.8
ViT-L	C100	0.61	0.47	0.6	2.1
DINO-B	IN	0.40	0.33	2.8	4.1
DINO-B	C100	0.60	0.45	0.0	1.4
DINOv2-S	IN	0.37	0.33	1.8	1.2
DINOv2-S	C100	0.60	0.44	0.0	0.8
DINOv2-B	IN	0.58	0.64	0.3	0.1
DINOv2-B	C100	0.66	0.54	0.0	0.1
DINOv2-L	IN	0.69	0.72	0.0	0.0
DINOv2-L	C100	0.70	0.64	0.0	0.1
DINOv2-G	IN	0.70	0.74	0.0	0.0
DINOv2-G	C100	0.77	0.70	0.0	0.0
CLIP-B	IN	0.36	0.29	4.2	7.9
CLIP-B	C100	0.60	0.46	0.0	3.1
CLIP-L	IN	0.35	0.27	4.0	8.4
CLIP-L	C100	0.57	0.43	0.0	2.2
SigLIP-B	IN	0.34	0.26	5.0	7.5
SigLIP-B	C100	0.59	0.43	0.0	2.3

Table 38: Edge-type ORC on the centroid kNN graph (30 WordNet superclasses): negative curvature concentrates on between-cluster bridges.

A.17 PROMPT ROBUSTNESS

model	photo	name	image	closeup	this is	wild	def.	gloss only	concept	discuss.
BGE-base	0.124	0.118	0.125	0.126	0.125	0.115	0.119	0.115	0.119	0.121
E5-base	0.114	0.119	0.117	0.120	0.115	0.116	0.108	0.115	0.120	0.116
GPT-2	0.131	0.145	0.110	0.086	0.080	0.089	0.105	0.124	0.112	0.105
GPT-2-M	0.091	0.122	0.110	0.100	0.134	0.104	0.081	0.107	0.112	0.098

Table 39: $\hat{\delta}$ across prompt templates (ImageNet class names; “gloss only” contains no class name).

template	GPT-2 S	GPT-2 M	GPT-2 L	GPT-2 XL
photo	−0.016*	−0.019*	−0.043*	−0.050*
name_only	−0.006	−0.007	−0.048*	−0.051*
image	−0.018*	−0.028*	−0.050*	−0.049*
closeup	+0.007	−0.018*	−0.044*	−0.049*
this_is	−0.032*	−0.022*	−0.042*	−0.046*
wild	−0.023*	−0.067*	−0.014	−0.010*
def_pair	−0.002	−0.019*	−0.036*	−0.035*
gloss_only	−0.020*	−0.030*	−0.033*	−0.036*
concept	−0.002	−0.003	−0.043*	−0.041*
discussion	−0.010	−0.013*	−0.041*	−0.040*

Table 40: Template-conditioned excess over the spectrum-matched null for the four GPT-2 sizes (10 templates; *: $|z| \geq 2$ against combined null and estimator noise). The scale trend of §4 is assessed on every template, not only on the paper’s baseline. Padding-free extraction (one prompt at a time, the canonical protocol of Table 14). Mean excess over templates: −0.012 (S), −0.023 (M), −0.040 (L), −0.041 (XL); cells at $|z| \geq 2$: 5, 8, 9 and 10 of 10. Only 3 null replicates per cell here, so the z are coarser than the census’s 20-replicate values: on the baseline template the 200-replicate census (Table 14) finds S and M not genuine (excess −0.013/−0.009, $p = 0.09/0.14$), so their starred cells are borderline, whereas the L/XL cells (z from −4 to −9) are not in doubt. Template sensitivity of raw $\hat{\delta}$: range up to 0.065 for causal LMs and ≤ 0.012 for embedders (prompt table in Appendix 39).

A.18 SAMPLE-LEVEL FEATURES

model	CIFAR-100 (10 images/class, $n=1000$)				DTD (22 images/class, $n=1034$)			
	sup. $\hat{\delta}$	$\hat{\delta}_{99.9}$	excess	$r/200$ (p)	sup. $\hat{\delta}$	$\hat{\delta}_{99.9}$	excess	$r/200$ (p)
ViT-T	0.130	0.093	-0.0016°	134 (0.333)	0.124	0.087	-0.0004°	111 (0.448)
ViT-S	0.121	0.086	$+0.0051^\circ$	21 (0.896)	0.103	0.070	-0.0062	197 (0.020)
ViT-B	0.121	0.082	$+0.0031^\circ$	50 (0.751)	0.092	0.060	-0.0080	200 (0.005)
ViT-L	0.100	0.067	-0.0041°	185 (0.080)	0.092	0.056	-0.0073	200 (0.005)
DINO-B	0.104	0.070	-0.0042°	194 (0.035)	0.110	0.075	$+0.0008^\circ$	80 (0.602)
DINOv2-S	0.100	0.065	-0.0049	200 (0.005)	0.145	0.090	$+0.0058^\circ$	0 (1.000)
DINOv2-B	0.089	0.050	-0.0090	200 (0.005)	0.139	0.073	-0.0019°	162 (0.194)
DINOv2-L	0.087	0.043	-0.0131	200 (0.005)	0.131	0.064	-0.0060	200 (0.005)
DINOv2-G	0.102	0.036	-0.0232	200 (0.005)	0.136	0.065	-0.0079	200 (0.005)
CLIP-B	0.117	0.080	-0.0060°	184 (0.085)	0.125	0.089	$+0.0021^\circ$	71 (0.647)
CLIP-L	0.132	0.084	$+0.0037^\circ$	35 (0.826)	0.120	0.082	-0.0008°	124 (0.383)
SigLIP-B	0.113	0.078	-0.0074	199 (0.010)	0.114	0.079	-0.0044°	184 (0.085)

Table 41: **The sample-level reading, the object of prior latent-hyperbolicity claims, sits within null noise in most cells.** The instrument on ≈ 1000 stratified training images per cell (the per-image features, not class centroids): sup. $\hat{\delta}$ is the raw supremum δ_{norm} the literature reports; $\hat{\delta}_{99.9}$, excess, $r/200$ and p follow Table 1 (Haar spectrum-matched null on the sample cloud, 99.9th-percentile statistic, 200 replicates); $^\circ$: not genuine under Benjamini–Hochberg over the 24 cells. The raw supremum lies in the band 0.087–0.145; the excess is not genuine in 14 of 24 cells and genuine in 10 (DINOv2-S, DINOv2-B, DINOv2-L, DINOv2-G, SigLIP-B on CIFAR-100; ViT-S, ViT-B, ViT-L, DINOv2-L, DINOv2-G on DTD), where it reaches at most -0.023 .

A.19 TEXT-MODALITY TRANSFER

model	supremum, Gaussian null		census of record			trip. (cos)	NC H–R (pp)	FS H–R (pp)
	$\hat{\delta}_{\max}$	excess	$\hat{\delta}_{99.9}$	excess	$r/200$ (p)			
BGE-base	0.122	-0.028	0.082	-0.021	200 (0.005)	0.89	-0.59	$+0.05 \pm 0.05$
E5-base	0.127	-0.024	0.086	-0.017	200 (0.005)	0.89	-0.43	$+0.04 \pm 0.03$
GTE-base	0.132	-0.020	0.088	-0.017	200 (0.005)	0.89	-0.82	$+0.08 \pm 0.06$

Table 42: DBpedia Classes (219 leaf classes, 3 levels): excess under the original supremum reading (3 Gaussian replicates) and under the census of record (Haar null, 99.9th-percentile statistic, 200 replicates; r = replicates above the real value, left-tail p), strong sibling alignment, and $\hat{\delta}$ above the low- δ band correctly predicting marginal gains. Tree map on the same classes: no configuration degenerates (largest cluster $\leq 10\%$), the embedders recover the true 70-way level-2 partition at ARI 0.21–0.29 in all 6 configurations and agree with each other at cross-model ARI 0.65–0.82.

A.20 HIERARCAPS MULTI-LEVEL HIERARCHY

On HierarCaps (Alper & Averbuch-Elor, 2024) (1000 four-level caption chains), text encoders show the same reading with no class set: the projection radius orders levels within a chain, sibling triplets are resolved at 0.83–0.97 (chance 0.5), and the leaf excess repeats the text census (Table 43).

model	$\rho(\text{level, radius})$	% monotone	trip. L1	trip. L2	$\hat{\delta}$ leaves	excess
BGE-base	+0.80 $\pm$ 0.31	49.6	0.88	0.97	0.130	+0.003
E5-base	+0.87 $\pm$ 0.19	53.9	0.90	0.97	0.139	+0.010
GTE-base	+0.47 $\pm$ 0.55	24.0	0.90	0.97	0.135	+0.006
CLIP-B	+0.64 $\pm$ 0.43	27.2	0.86	0.95	0.123	-0.011
CLIP-L	+0.51 $\pm$ 0.49	18.7	0.86	0.96	0.117	-0.010
SigLIP-B	+0.32 $\pm$ 0.55	9.8	0.83	0.91	0.138	+0.014

Table 43: HierarCaps test set (1000 four-level chains): radial ordering under the paper’s projection (chance for strict monotonicity = $1/24 \approx 4.2\%$; shuffle controls ≈ 0), sibling triplet agreement (cosine; chance 0.5), and leaf $\hat{\delta}$ with spectrum-null excess.

A.21 PARALLELOGRAM DEFECT ξ : A NORM-STRUCTURE DESCRIPTOR

ξ reads norm structure, not curvature. A parallelogram estimator ξ , intended as a curvature sign, is driven negative by heterogeneous norms alone, so it is reported only as a norm-structure descriptor. For a triangle (a, b, c) with m the approximate midpoint of (b, c) (the minimizer of $D(\cdot, b) + D(\cdot, c) + |D(\cdot, b) - D(\cdot, c)|$), the parallelogram defect is $\xi = (D(a, m)^2 + \frac{1}{4}D(b, c)^2 - \frac{1}{2}D(a, b)^2 - \frac{1}{2}D(a, c)^2) / (2 D(a, m) \text{diam})$, averaged over 2×10^4 sampled triangles: 0 in Euclidean space, positive on spheres, negative in hyperbolic space and on trees.

space / model	ξ (mean)	frac. negative	null ξ	excess ($\pm$ CI95)	z
iid Gaussian ($d=768$)	+0.092	0%	—	—	—
S^d (geodesic)	+0.113	0%	—	—	—
S^d (chord)	+0.115	0%	—	—	—
$\mathbb{H}^2$ region ($R=6$)	-0.194	96%	—	—	—
balanced tree (depth 9)	-0.239	96%	—	—	—
ViT-T	+0.027	25%	+0.027 $\pm$ 0.002	-0.001 $\pm$ 0.003	-0
ViT-S	+0.027	23%	+0.042 $\pm$ 0.002	-0.015 $\pm$ 0.005	-7
ViT-B	+0.033	16%	+0.047 $\pm$ 0.002	-0.014 $\pm$ 0.004	-6
ViT-L	+0.028	21%	+0.046 $\pm$ 0.002	-0.018 $\pm$ 0.005	-7
DINO-B	-0.008	55%	+0.052 $\pm$ 0.003	-0.059 $\pm$ 0.005	-21
DINOv2-S	-0.020	66%	+0.061 $\pm$ 0.004	-0.081 $\pm$ 0.008	-20
DINOv2-B	-0.049	83%	+0.075 $\pm$ 0.003	-0.123 $\pm$ 0.005	-45
DINOv2-L	-0.068	90%	+0.080 $\pm$ 0.002	-0.149 $\pm$ 0.004	-66
DINOv2-G	-0.070	89%	+0.082 $\pm$ 0.002	-0.152 $\pm$ 0.004	-69
CLIP-B	+0.019	32%	+0.024 $\pm$ 0.002	-0.005 $\pm$ 0.003	-3
CLIP-L	+0.028	28%	+0.032 $\pm$ 0.002	-0.004 $\pm$ 0.004	-2
SigLIP-B	+0.036	20%	+0.031 $\pm$ 0.002	+0.005 $\pm$ 0.004	+2

Table 44: Parallelogram curvature-sign estimator ξ on reference geometries (top) and ImageNet centroids (bottom), with the spectrum-matched null, excess and z -score for each model (20 null replicates; noise combines null and estimator s.d.): matched nulls have positive ξ , so raw readings understate curvature.

model	ξ_{geo}	frac. neg.	excess	z
ViT-T	+0.031	0.24	−0.0048	−2.7
ViT-S	+0.065	0.08	+0.0092	+6.2
ViT-B	+0.067	0.07	+0.0144	+10.5
ViT-L	+0.074	0.04	+0.0113	+7.9
DINO-B	+0.085	0.05	+0.0014	+1.0
DINOv2-S	+0.100	0.02	+0.0033	+2.3
DINOv2-B	+0.108	0.01	+0.0046	+4.9
DINOv2-L	+0.108	0.01	+0.0040	+3.9
DINOv2-G	+0.094	0.01	−0.0114	−12.5
CLIP-B	+0.025	0.28	−0.0129	−7.1
CLIP-L	+0.033	0.25	−0.0120	−4.8
SigLIP-B	+0.051	0.14	+0.0025	+1.2

Table 45: ξ on the angular geometry: centroids L2-normalized, distances the spherical geodesic, spectrum null built on the normalized cloud and re-normalized (20 replicates; ImageNet centroid store). Probes whether the deep negative ξ -excess of DINOv2 on ImageNet (Table 44) is a norm-structure effect. Outcome: it is for S/B/L, whose angular ξ -excess turns sphere-leaning positive, while DINOv2-G and CLIP-B/L remain negative on the sphere: the Euclidean ξ signal of DINOv2 is largely carried by norm structure.

reference (no curvature unless stated)	ξ	frac. negative
iid Gaussian	+0.0919	0.00
Gaussian, lognormal radii ($\sigma=0.3$)	−0.0621	0.99
Gaussian, lognormal radii ($\sigma=0.6$)	−0.0607	1.00
Gaussian, lognormal radii ($\sigma=1.0$)	−0.0496	0.99
core + shell mixture	−0.3217	0.79
uniform sphere	+0.1154	0.00
sphere with radial jitter	−0.0571	1.00

Table 46: ξ on references with heterogeneous norms but no curvature ($n=1000$, $d=768$, 5 seeds). Norm dispersion alone drives ξ negative (lognormal radii, radial jitter on a sphere, core+shell): the estimator reads norm structure, not curvature, and the paper reports it only as such.

B ILLUSTRATIVE VISUALIZATION

Figure 12 is a low-dimensional illustration of the hub-and-spoke organization; every quantitative comparison in the paper is made in the full feature space (§4).




Figure 12: Illustrative 2-D projection of the 100 CIFAR-100 class centroids (per-panel PCA fit on the 20 super-class hubs; panels preserve 22–30% of hub variance and 6–10% of spoke length, so visual compactness is not comparable across panels and quantitative comparisons are made in full dimension). (a) DINOv2-S, Euclidean; (b) DINOv2-S after the projection of §3; (c) DINOv2-G.